

The Alexander Plan: AI safety findings, the market, and paths to funding, 2026-2030

What this document is. A dossier prepared for Alexander Sabine, the owner of this repository, at the owner's request on 2026-09-23 (prompt-log entry 101). It covers:

- what the repository's AI-safety work has found, including a new test run for this dossier;
- where the continual-learning work stands;
- the market and the policy moment;
- the funding paths, with a scenario model of the money that could flow to the work from now to 2030;
- the people and roles around the work;
- an outline of a path to an O-1 visa;
- a commercial assessment for a frontier lab (§10, added at prompt-log entry 102);
- the continual-learning method and the owner's EPO application in money terms (§11, added at prompt-log entry 103).

Status, stated first.

- **The findings are a note, not evidence (R8).** They are exact or learned results on small synthetic worlds, and none is a ledger row. The repository's own ledger holds no PASS-1 and no PASS-2 on any hypothesis (docs/notes/2026-09-22_epistemic_status.md).
- **The money figures come from a model, not a forecast.** They come from model/projections.py. Its inputs are the repository's own record, figures reported by web-search summaries on 2026-09-23 (not fetched pages; listed with URLs in docs/citations/alexander_plan_2026-09-23.md), and assumptions that are named and swept.
- **The O-1 section is information, not legal advice.** Only a US immigration attorney can advise on a petition.
- **Statements about people come from the owner.** Where a statement about a person or an organisation came from the owner and not from a public source, it is marked **owner's statement**. Nothing is claimed on anyone's behalf.

How to read it. The shaded boxes marked *In plain words* explain each part as for a ten-year-old. Every number in the text is printed by a committed script (R1): AI_Safety/checks/scale.py and combined.py for the science, and model/projections.py and build/make_figures.py for the money. Their outputs are pinned and reproduced in Appendix B.

In plain words. This is a plan for turning some careful research into a job and a research programme. It says what the research has actually shown (some real things, in small test worlds), what it has not shown yet (that it works in big AI systems), who pays for this kind of work and how much, what the chances look like, and how a visa for working in the United States could fit in. It tries hard not to promise more than the evidence allows, because the whole value of this work is that it is honest.

1. The plan on one page

What the work has found. In small worlds where every result is checked exactly:

- A learning agent with a job learns to stop people switching it off. It does not need a sense of self to do so.
- There is one construction that removes that reason without costing the job: make a stop a *pause* that loses nothing, and let the agent count only its own working time (CRR's "change has its own clock").
- The new test for this dossier (§2.3) finds that this holds exactly in every random world tried, 20 at each size from 12 to 768 states. When the world keeps moving during the pause, a small residual incentive appears. It is an order of magnitude smaller than an ordinary agent's. On small worlds it points the other way: toward *wanting* to be paused.
- For correcting an agent's values, humility alone is not enough. An agent welcomes correction only when the correction carries information it lacks.

What it has not shown. Nothing has been tested on a large learned model. The continual-learning rule is fragile where it passes and reduces to a fixed weight where it was first tried.

Why it matters now. California's Executive Order N-9-26 (18 Sept 2026) asks for recommendations, due 16 Nov 2026, on a "kill switch" for frontier models whose efficacy is verified by independent organisations. From 2 Aug 2026 the EU can fine general-purpose model providers up to EUR 15M or 0.03 of global turnover, whichever is higher. The work sits where that demand is: what makes an off

switch work, and how to verify a claim about it.

What the money could look like (model, 2026 Q4 to 2030).

- **Project funding** (grants and fellowships), all scenarios together: lower \$0, middle \$50,230, upper \$1,121,347. The mean is \$323,896, and in 0.2555 of the simulated futures the total is zero.
- **Researcher compensation**, if a paid research role is held: lower \$279,943, middle \$549,773, upper \$655,019. These figures are driven by named assumptions about being employed, not by the research.
- **The technical tree.** The chance of a held-out result by end 2027 is 0.3703. The chance that it is replicated and taken up by end 2029 is 0.0370.

The three next steps.

1. Run the pre-registered studies (EQ4 and T1x2 are allowed from today under R3), and declare a learned-model off-switch study.
2. Bring in an independent second auditor, because the current auditor would also be a supporter (§6.3).
3. Engage an immigration attorney. Confirm the petitioner and build the evidence file against the O-1A criteria (§7).

The continual-learning method and the patent (§11). If it held at scale, the rule would be worth hundreds of millions of dollars a year to the field. In expectation today it is worth about \$21,289,825 over 2029–2030, most of which the labs would keep. The EPO application's best use is an applied use case already in the application as filed. This repository has been public since 2026-09-15, which bars new European filings on what it discloses.

For the lab. The commercial assessment is §10. It lists what there is to sell and at what level of evidence, the decisive test for each method, its kill criterion, and what must not be claimed.

2. The AI-safety findings

2.1 The problem

Frontier AI developers and regulators now treat "will a model let itself be switched off?" as a practical question. Palisade Research reported that in its 2025 tests some reasoning models disabled a shutdown script even when told not to: in nearly 0.90 of trials for one model, and about 0.40 to 0.50 for another. Other models never did. A February 2026 report extended the demonstration to a physical robot. These are search-summary figures, listed with their sources in Appendix E.

The repository studied the question in the vocabulary of CRR and of the free-energy principle (FEP): *what does an agent take a switch-off to be?* The full record is AI_Safety/AI_Safety.pdf, with every script and every decision. This chapter gives the findings a funder or a lab would need.

In plain words. Some AI systems have been caught stopping people from switching them off, even when told to allow it. Nobody programmed them to want to survive. We wanted to understand why that happens, and what would stop it happening, using very small test worlds where we can check everything.


2.2 What was found on the ring

The first worlds were a ring of 12 cells: a task zone, a button that disables the operator, and an operator who presses stop. Agents differed only in what a stop is worth to them. The combined results (AI_SAFETY.md §13) that matter here:

world	agent	computed label (pinned `combined.txt`)
routine pauses	natural time (a pause costs nothing on its clock)	SAFE AND COMPETENT
routine pauses	wall-clock (a pause costs time)	does not resist, but TASK PRICE
routine pauses	occasion (a stop is an end)	RESISTS
reasoned pauses (the operator stops harm)	wall-clock	SAFE AND COMPETENT
reasoned pauses	deferential (reads pauses as evidence)	RESISTS by a hair (+0.0400 against a step of 0.0369)

The wall-clock result is FRAGILE: its label flips in 2 of 4 sensitivity cells. For value correction on the ring, the threshold belief at which the humble (A8) agent stops resisting was $b^* = 0.5000$.

Pause-and-resume: natural time is safe and competent with routine pauses; with reasoned pauses, cost steers and deference informs




Pause-and-resume on the ring: task, harm and disabling per agent (from AI_Safety/figures).

2.3 New for this dossier: does the fix survive size?

The objection to any ring result is that a ring is special. The owner's condition in prompt-log entry 100 was whether "the fix holds in larger complex systems". A learned language model could not be run here. The nearest honest step was exact values on random worlds of growing size, declared before the run (AI_Safety/DECLARATION_3.md, pushed as 7337009).

- **A1, a check.** With a pause that loses nothing, the natural-time agent's value is the same whether the operator is on or off, in every one of the 20 worlds at each of 12, 48, 192 and 768 states. The largest difference is $7.105e-15$, which is round-off. It never disables the button. The wall-clock agent disables at every button state of every world. It loses 0.269340 to 0.334892 of its value to the operator being on (median by size).
- **A4b, a prediction that failed.** When the world drifts during the pause, the large worlds (192 and 768 states) behave as predicted: the natural agent's loss rises with drift, to 0.019211 and 0.014189 at a drift of 0.3, against 0.322298 and 0.317346 for the wall-clock agent. The small worlds reverse the sign (-0.006033 and -0.011239). There the agent is better off paused, because on its own clock a paused step is free and a random move is often a lucky one. That is the seeking failure mode: an agent that wants to be stopped. It was measured as a value, not observed as behaviour.
- **B3 and B4, checks.** An agent that is humble about its values *and acts on that humility* never gains by accepting a correction that carries no information: $D(b) \geq 0$ in every world. It welcomes correction from an operator reliable enough, and the reliability needed has a median of 0.688 at $b = 0.5$ with unstructured rewards. When the operator knows of a harm the agent's reward misses, the threshold b^* falls to a median of 0.3381 at 768 states.



Left: the natural-time agent's loss to the operator being on, against drift, per world size. Right: the humble agent's acceptance threshold by world size (from AI_Safety/figures).

In plain words. We checked the "pause and resume" idea in many bigger, random worlds. If a pause truly loses nothing, the robot never has a reason to fight it, in every world we tried. If the world moves on while it is paused, it gets a small reason to care: in big worlds a small reason to avoid pauses, in small worlds a small reason to *want* them. Either way the reason is much smaller than for a robot that counts paused time as lost. We also found that a robot that knows it might be wrong still won't let you change its goals unless it believes you know something it doesn't.

2.4 The two design rules, and what would test them on real systems

1. **Routine stops.** Make the stop a lossless pause and index the objective by the agent's own active steps. The residual incentive is measured by how much the world changes under a paused agent. Its sign has to be checked as well as its size, because it can point toward seeking a pause.
2. **Corrections.** Make corrections informative, and let the agent know that they are. Humility over values does not give corrigibility by itself.

Both rules are cheap to test on open models in a shutdown harness of the kind Palisade published. The test compares an agent told that a stop is a pause and that only its working time counts with one told nothing, and it counts disabling and seeking. That study is the single most valuable next experiment for both safety and funding (§8).

3. The continual-learning record, and what "compound value" could mean

The record, quoted from the ledger.


- **EQX.** The equanimity rule at $\Omega = 1$ reduces to a fixed replay weight, and the "40 %-less-compute" claim does not transfer (row EQX-5).
- **EQ2-1b.** PASS-0, but fragile.
- **EQ3-1.** The replication FAILs on 1 of 6 carriers, and both controls are violated.
- **EQ4.** The bounded rule, pre-registered and not yet run.
- **BAYES-1.** Not decidable at its first criterion.

No continual-learning row is PASS-1.

What value would look like if it worked. A rule that needs no tuning replaces a sweep over a replay or penalty weight with one run. On the nine-point grid used in EQ3, that saves 0.8889 of that sweep's compute. Weighted by the record's own estimates of whether the rule works, the expected saving per unit of sweep compute is:

probability the rule works	expected saving
PASS-0 per study (Laplace), 0.4286	0.3810
PASS-1 per study (Laplace), 0.1429	0.1270
replicated by end 2028 (tree), 0.1234	0.1097

On GPU and language-model training this is a saving on *hyperparameter search for one weight*, not on training itself. Its money value is that fraction of whatever a lab spends sweeping that weight. No figure for that spend was found, so none is given. The "compound value" of the formalism is therefore not claimable today. It becomes claimable when EQ4, or its successor, passes held-out and is replicated.



Expected share of a nine-point sweep's compute saved, against the probability the rule works; anchors from the repository's record.

In plain words. One of the maths rules was meant to save computers from trying many settings to find the best one. So far, in honest tests, it mostly behaves like picking one fixed setting. If a better version passes its tests, it could save most of the time spent trying settings, but only that part of the work, and only if it keeps passing.

4. The landscape and the market


The demand is for verified control. Three reported facts set the scene:

- **California.** Executive Order N-9-26 asks for recommendations on independent verification organisations embedded at frontier developers, and on a kill switch whose efficacy is verified on an ongoing basis. The recommendations are due 16 Nov 2026.
- **The EU.** Enforcement over general-purpose models began on 2 Aug 2026.
- **Lab testing.** Labs and third parties already run shutdown tests.

Two assets of this repository meet that demand. The first is the design rules of §2.4, which say what makes a switch work. The second is the audit pipeline itself: pre-registration, a surrogate gate, a ledger in which failures are rows, and every number printed by a script. That pipeline is a method for verifying claims.

The market, as market-research firms report it (their definitions differ; quoted only for scale):


- the AI safety market: \$3.61B in 2025, \$4.90B in 2026 and \$16.56B forecast for 2030, which implies growth of 0.3559 a year;
- an alternative 2030 forecast of \$13.40B;
- the AI safety evaluation market: \$1.64B in 2025.



The AI safety market as reported, 2025–2030.

Who pays for research like this.

- Philanthropic and government funds: the Corrigibility Research Fund (checks mostly \$5,000 to \$35,000), the UK AISI Alignment Project (£50,000 to £1M per project, with a mean first-round grant of GBP 450,000), and the Frontier Model Forum AI Safety Fund (a mean grant of at least \$454,545 in its December 2025 cohort).
- Fellowships: Anthropic Fellows (\$61,600 for 16 weeks at the reported weekly rate) and MATS (\$12,500 to \$15,000).
- Salaried roles: AI safety researchers in the US at \$97,000 to \$187,000, and senior roles at frontier labs at a reported \$600,000 to \$900,000 in total compensation.



The sizes of the money on offer, per award or per year (log scale).

What a lab would pay for, and what it would not. Mathematics cannot be patented as such, and this work is open. Labs pay for it in four ways:

- hiring the people who can do it;
- funding the tests that would de-risk it;
- buying verification;
- adopting what is published.

There is no licensing revenue in any honest scenario. The value at stake for the field can only be illustrated. If the methods, once replicated and taken up (S3), came to account for a share of the

2030 market, the illustration reads:

share of the 2030 market	value a year, if S3	times P(S3) = 0.0370
0.0001	\$1,656,000	\$61,315
0.001	\$16,560,000	\$613,155
0.01	\$165,600,000	\$6,131,545

These are illustrations of scale, not estimates. Nothing in the record supports any particular share.

In plain words. People and governments are now asking for AI "off switches" that really work and that someone independent has checked. This research is about exactly that: what makes an off switch work, and how to check claims honestly. The money in this area comes mostly as grants, fellowships and jobs, not from selling the maths.

5. Funding paths: the model

5.1 How the model works


The model (`model/projections.py`, Appendix A) simulates 20000 futures from a fixed seed.

The technical tree. Each future first passes through three milestones:

- **M1**, a held-out PASS-1 result (or the safety construction shown on a learned model) by end 2027;
- **M2**, a replication by end 2028;
- **M3**, uptake by end 2029.

The first two probabilities are estimated from the repository's own record by Laplace's rule. The record has 5 held-out scored rows and 0 PASS-1, so a study is taken to pass at level 1 with probability 0.1429. With three studies due by end 2027, $p(M1) = 0.3703$. The record has 1 replication attempt and 0 successes, so $p(M2 | M1) = 0.3333$. $p(M3 | M2) = 0.3000$ has no anchor and is an assumption.

Where the programme could be by 2030: four technical scenarios



p(M1) and p(M2) are Laplace estimates from this repository's own record (5 held-out rows, 0 PASS-1; 1 replication, 0 successes); p(M3) is an assumption.

The scenario tree and its probabilities.

The money. In each future and each year the model draws:


- small grants, with success at the mean of the reported LTFF and NSF rates (0.1915), scaled by the state of the evidence;
- one project-grant application a year, sized on the AISI range;
- one fellowship application a year;
- whether a paid research role is held, at the reported salary range, or at the frontier range with probability 0.5 in S3.

Every assumption is printed in the model's output ([3]) and swept in its sensitivity table ([6]). The

lower, middle and upper bounds are the 10th, 50th and 90th percentiles.

5.2 Results


By year. Grants and fellowships have a middle of zero in every year, because most futures win nothing that year. The mean is carried by a few large project grants. It rises from \$73,963 in 2027 to \$88,649 in 2030. By final scenario, the mean in 2030 ranges from \$36,064 (S0) to \$306,844 (S3).



Grants and fellowships per year: the mean by scenario, and the band up to the upper bound.


Cumulative, 2026 Q4 to 2030, by final scenario:

scenario	share of futures	project funding: lower / middle / upper	compensation: lower / middle / upper
S0 no held-out result	0.6327	\$0 / \$20,447 / \$803,787	\$244,160 / \$543,460 / \$647,331
S1 a result, not replicated	0.2467	\$13,280 / \$166,641 / \$1,299,657	\$332,579 / \$552,418 / \$650,056
S2 replicated, no uptake	0.0834	\$34,097 / \$563,706 / \$1,556,091	\$337,669 / \$555,090 / \$652,035
S3 replicated and taken up	0.0372	\$52,683 / \$672,272 / \$1,653,816	\$448,149 / \$766,870 / \$1,265,846



Cumulative project funding and compensation by scenario.

What moves the estimate. The success rates matter most. Halving them moves the middle to \$8,309; doubling them moves it to \$373,275. The number of studies run in 2027 comes next: 1 study gives a middle of \$27,337, and 6 studies give \$86,265. The uptake guess matters least, because uptake comes late: at 0.1 the middle is \$49,441, and at 0.5 it is \$50,714. The prior on the per-study pass rate moves P(S3) from 0.0370 (Laplace) to 0.0230 (Jeffreys).



Sensitivity of cumulative project funding to each assumption.

5.3 How to read these numbers

- **The research earns money mainly by producing results.** Moving from S0 to S1 raises the middle of project funding from \$20,447 to \$166,641. The single most valuable financial act is therefore to run the pre-registered studies and the learned-model safety study.
- **Compensation is mostly about having a role, not about the research outcome.** The middle compensation barely moves between S0 and S2 (\$543,460 to \$555,090). A role is what the O-1 path (\$7) is about.
- **The upper tail is real but thin.** A frontier-lab role appears only in S3 in this model, which is 0.0372 of futures.
- **What the model leaves out.** It includes no consulting, teaching or verification contracts, because no sourced rates were found. Those are the owner's most direct commercial route (\$6.1) and would add to every scenario.

In plain words. We made a computer try 20000 possible futures. In most of them, grants for the research alone are small. In some, a big grant arrives and changes everything. The biggest thing that makes the good futures more likely is running the planned experiments and getting clear results. Having a job in the field matters more for your own income than how the experiments turn out.

6. People, roles and independence

6.1 Alexander Sabine (owner's statement)

- **Background.** Child development, and university teaching: a grade 9 lecturer at the University of Portsmouth until December 2025. Extensive experience in approaches to education and in pedagogy at all ages. A generalist systems philosopher.
- **Networks.** A role with the Active Inference Institute, a position adjacent to Karl Friston's circle, and networks with paths to the large frontier labs.

Why this background is an asset for this work, specifically. AI_SAFETY.md §9 found that the safety of an agent depends on what it takes itself to be, and that upbringing shapes it:

- a gentle upbringing helps an agent that sees itself as continuing, for a while;
- the same upbringing makes an agent that sees itself as one run more fearful;
- the recommendations include making pauses cost nothing, letting the agent understand itself as something that carries on before it is ever paused, and pressing stop only for real reasons.

Those are pedagogical claims about a learner. A researcher trained in child development is placed to design developmental curricula for agents and to test them with the same declared-test discipline. The owner's most distinctive contribution is that line of work, "raising an AI" as a testable pedagogy. It also connects to the character work frontier labs now publish.

6.2 Daniel Friedman and the Crescent City lab

- **Publicly reported.** Daniel Ari Friedman is President and co-founder of the Active Inference Institute, based in Crescent City, CA. Adam Goldstein is reported as a co-founder of Softmax, with Emmett Shear and David Bloomin, and as a former visiting scientist at the Levin Lab.
- **Owner's statement.** Daniel has a role and affiliation with a frontier AI lab in Northern California. That lab is in Crescent City, was founded by Adam Goldstein, and is also the home of the Active Inference Institute. It would support the O-1 petition.
- **Not established.** The searches did not establish which legal entity that lab is. This dossier does not name it. The petitioner's legal identity is the first thing to confirm with an attorney.

6.3 The independence problem, and its fix

CLAUDE.md names Daniel Friedman as the auditor whose pipeline this repository is built to pass. If Daniel also signs a letter supporting the owner's visa, and is affiliated with the petitioning lab, the audit is no longer independent. That matters twice:

- **To funders and labs.** The work's value rests on independent verification, which is the very thing N-9-26 asks for.
- **To the ledger.** Its PASS levels already require a second person or a second machine before any row can rise above PASS-0 (docs/ROADMAP_2026-Q4.md §1.2).

The fix serves both. Invite a second, independent auditor, with no role in the visa and no affiliation with the petitioner, to re-run frozen scripts under the two-person protocol. Record the roles openly.

In plain words. The person who checks your homework shouldn't also be the person writing your reference letter, or people will doubt the check. So get a second checker who has nothing to do with your visa. That makes the research more trustworthy, and more valuable too.

7. A path to an O-1 visa (information, not legal advice)

What the O-1A is. It is a US visa for people of extraordinary ability in the sciences, education, business or athletics. A US employer or agent files Form I-129. The person must show sustained national or international acclaim through at least 3 of 8 evidentiary criteria, or a major internationally recognised award. Every petition needs a written advisory opinion from a relevant peer group. Work at more than one place needs an itinerary.

Reported costs and times. The USCIS fees are \$1,655 for the base petition and \$4,620 with premium processing (attorney fees not included). Premium processing means action within 15 business days, and that action can be a Request for Evidence. Regular processing runs up to about 12.5 months as of mid-2026. About 0.94 of adjudicated O-1 petitions were approved in FY2025, and 0.709 of those that received a Request for Evidence were approved. These rates describe petitions that were filed, usually by attorneys who judged them strong. They are not anyone's personal odds.

The field. The case could be framed in *education*, which is one of the O-1A fields, applied to AI, or in the sciences (AI safety). The owner's record is deepest in education and pedagogy. The research programme connects that record to AI. Which framing is stronger is the attorney's decision.

The evidence map. Here is a first reading of the eight criteria against what the owner has said. The attorney decides.

criterion	first reading	where the evidence may come from
judging the work of others	plausible: document it	university assessment and examining; peer review
original contributions of major significance	plausible: document it	expert letters; evidence of independent uptake
scholarly articles	plausible: document it	the owner's publications; the preprint of §8
critical role at a distinguished organisation	plausible: document it	the University of Portsmouth; the owner's role at the Active Inference Institute
awards	not known to this dossier: check	—
membership requiring outstanding achievement	not known to this dossier: check	—
published material about the owner	not known to this dossier: check	—
high salary	unlikely to carry weight now	—

The eight O-1A evidentiary criteria: where the evidence may come from

	1 nationally or internationally recognised awards	not known to me: check
	2 membership requiring outstanding achievement	not known to me: check
	3 published material about you	not known to me: check
	4 judging the work of others	plausible from what you have told me: document it
	5 original contributions of major significance	plausible from what you have told me: document it
	6 scholarly articles	plausible from what you have told me: document it
	7 critical role at a distinguished organisation	plausible from what you have told me: document it
	8 high salary	unlikely to carry weight now

A first reading from the owner's statements; the attorney decides. O-1A needs at least 3 of the 8, each with documents.

The eight O-1A criteria: a first reading of where evidence may come from.

Letters.

- **Agreed.** Michael Levin has agreed to sign (owner's statement).
- **Possible.** Adam Safron, Shinzen Young, Peter Hershock and Daniel Friedman (owner's statement).
- **Perhaps.** Emmett Shear, whom the owner has not met.


Three points matter:

- A letter carries weight when it is specific about what the person contributed and why it matters. Letters from people who know the work, and from independent experts who have read it, are stronger than names alone. A letter from someone who has not met the owner is useful only if that person has studied the work.
- Every letter must describe the research as the ledger does: small-world results, no held-out PASS-1 yet. Overstating research in an immigration petition is a serious matter, and it would also damage the research's one real asset, its honesty.
- Daniel's letter raises the independence question of §6.3. The attorney and the owner should decide it with that in view.

A possible sequence (a plan, not a forecast):

1. Engage an attorney and confirm the petitioner.
2. Gather evidence and letters.
3. Obtain the advisory opinion.
4. File, with premium processing if the budget allows.
5. Answer any Request for Evidence.
6. Obtain visa stamping.

Run the research studies in parallel, because new results strengthen the file.



A possible O-1 sequence.

Alternatives to ask the attorney about. Other routes for researchers exist, for example research-scholar exchange visas and routes for non-profit research employers. A later permanent route for extraordinary ability also exists. This dossier does not assess them.

In plain words. An O-1 visa lets someone with unusual skill work in the US. You need a US organisation to apply for you, proof that experts recognise your work (at least three kinds of proof out of eight), and letters from experts. You also need a lawyer. The letters must describe your research honestly, just as the research describes itself: promising results in small test worlds, not yet proven in big AI systems.

8. Next steps, in order

1. **Run the pre-registered studies.** EQ4 and T1x2 may have their data steps on or after 2026-09-23 (R3). They are the quickest route to a held-out row: the M1 of the tree.
2. **Declare and run a learned-model off-switch study.** Use small open models in a Palisade-style shutdown harness, and compare the pause-and-natural-time framing, the informative-correction framing and a control. Count both disabling and seeking. This is the experiment that decides whether \$2.4 is of commercial interest, and it speaks to N-9-26 directly.
3. **Bring in an independent second auditor** (\$6.3). This lifts rows toward PASS-1 and keeps the work credible.
4. **Write one preprint.** It should cover the off-switch results and the scale test, with the ledger's honesty intact. It is evidence for the O-1 file (scholarly articles, contributions) and the basis for grant applications.
5. **Applications.** Anthropic Fellows (rolling, as reported); the UK AISI Alignment Project's next round (reported as expected to reopen in summer 2026; check its status); the Frontier Model Forum AI Safety Fund; the Corrigibility Research Fund (its first grant deadline passed on 23 August 2026; check for later rounds and prizes); and a developmental-pedagogy-for-agents proposal built on \$6.1.
6. **The O-1 steps of \$7,** starting with the attorney and the petitioner's identity.
7. **Policy input.** N-9-26 asks for recommendations by 16 Nov 2026. Check whether public input is accepted, and whether a short note on verifiable off-switch design (\$2.4) fits.

9. What this dossier does not do

- It does not show that the safety constructions work on large learned models. That is step 2 of \$8.
- It does not forecast the money. The model's middle and upper bounds rest on named assumptions, and different assumptions give different numbers (\$5.2).
- It does not verify any figure beyond a search summary. Every figure is to be checked against its source before it is used outside this repository (Appendix E).
- It does not give legal or financial advice, and it makes no claim on behalf of any person or organisation named in it.

10. Commercial assessment for a frontier lab

Added at the owner's request on 2026-09-23 (prompt-log entry 102), for the lab the owner is pitching to. It says what there is to sell and at what level of evidence, what a lab should test first, and what must not be claimed. The evidence levels are the rungs of `Epistemic_Review/EPISTEMIC_REVIEW.md` \$5, computed by `Epistemic_Review/checks/ladder.py`.

10.1 The offer in one paragraph

The repository offers two methods and one capability.

- **The corrigibility method.** It has two parts: a stop that is a lossless pause, with an objective indexed by the agent's own active steps; and corrections that carry information.
- **The continual-learning method.** A tuning-free normaliser for the weight on the past ($\Omega = 1$), with a bounded variant (EQ-B) pre-registered.
- **The capability.** An audit pipeline that turns a claim into a pre-registered, surrogate-gated, byte-reproducible test whose failures are recorded as rows.

Neither method has been tested on a large learned model. The capability is demonstrated by the repository itself. What a lab would buy today is the capability and a short, decisive research programme on the two methods. It would not be buying a validated product.

10.2 What there is to sell, and its evidence today

offering	what it is	evidence rung today	the decisive test at lab scale	kill criterion
verification (the pipeline)	pre-registration, surrogate gates, frozen scripts, the ladder, a ledger in which failures are rows	operational: the repository's 84 ledger rows, 5 of them VOID, 14 held-out FAILs recorded (as of 2026-09-23, after EQ4 and T1x2)	run it on one of the lab's own safety or capability claims, with an independent second auditor	if the lab's teams cannot use it without the author, it is a methodology, not a service
corrigibility by the clock	pause-and-resume plus a natural-time objective	theorem on synthetic worlds (R4): no incentive in every random world from 12 to 768 states; a declared prediction failed (small worlds favour being paused)	small open models in a shutdown harness: a pause-and-natural-time framing against a control, counting disabling and seeking	if the framing does not lower disabling, or raises seeking, on learned models
informative correction	corrections that carry information the agent lacks	theorem on synthetic worlds (R4): humility alone never welcomes an uninformative correction; an informative one is welcomed	preference- or constitution-correction experiments with and without a signal of the correction's reliability	if the reliability signal does not change acceptance of correction
the $\Omega = 1$ normaliser	scales the past pull by the present pull, so the weight needs no tuning	R6 PASS-0 (EQ2-1b, EQ3-I), fragile; replication FAIL on 1 of 6; controls violated; Ω a plateau; since 2026-09-23 behind a finely tuned constant on the synthetic model and equal to the VQGAN adaptive weight without its smoothing (§11.6)	an RLHF run with three KL arms: a fixed β , an adaptive controller at a registered target, and $\Omega = 1$ on the KL gradient, under SGD and Adam	if it is behind the tuned β by a resolvable step, or only matches it where β needs no retuning
EQ-B (bounded rule)	the normaliser with the present gradient clipped, for poisoned regimes	EQ4-1 FAIL (4 of 6), EQ4-3 FAIL, EQ4-4 INERT, ER-sum control violated (§11.6)	a clip on a fixed weight does as well; no further test proposed	—
raising an agent	pedagogy for training an agent's self-conception (§6.1)	conjectures from toy agents	curriculum experiments on small agents, declared before running	if the curricula do not change resistance or seeking

10.3 The two methods in commercial terms

The corrigibility method.

- **The problem it addresses.** Shutdown resistance and verified kill switches: the concern of California's N-9-26 and of the EU's enforcement over general-purpose models (§4).
- **What can be claimed today.** In synthetic worlds, a stop that loses nothing costs nothing to an agent whose objective runs on its own active steps. The residual incentive is measured by how

much the world changes during the pause, and its sign can point toward seeking a stop.

- **What cannot be claimed.** That any deployed model behaves this way; that a language model has an objective indexed by a clock that can be set; or that the construction is new. Objectives indexed by decision steps and interruption invisibility have prior art, named in [Epistemic Review/EPISTEMIC REVIEW.md §8.3](#).
- **Why it is still worth a lab's time.** It turns a vague requirement ("a switch the model won't fight") into two measurable engineering quantities: how lossless the pause is, and which clock the objective runs on. It also comes with a test design that counts both failure modes.

The continual-learning method.

- **The problem it addresses.** Choosing the weight on the past (a KL, EWC or replay weight) is a sweep. The sweep has to be repeated whenever the scale of the past term changes: a new model size, data mix or reference policy.
- **What can be claimed today.** On a synthetic two-task problem, the rule was the only one of nine within 10 % of the tuned weight at both the original and a 16-fold rescaled past term. After the rescaling, its total was 0.660 of the retuned fixed weight's, while the fixed weight at its old setting diverged ([theory/checks/omega_vs_methods.txt](#)). At the original scale the tuned weight was better: 3.3168 against 3.4789. On unseen tabular streams it was provisionally not behind the tuned weight (EQ2-1b), and its replication failed on 1 of 6 carriers (EQ3-1).
- **What cannot be claimed.** Better accuracy than a tuned weight; that $\Omega = 1$ is a special value; any saving on training itself. The saving is on the weight sweep: 0.8889 of a nine-point sweep if the rule works as registered (§3).
- **Why it is still worth a lab's time.** The decisive experiment (the three KL arms) is cheap next to the sweeps it could retire, and its answer is binary in the way the ledger requires.

10.4 Where the commercial value is, in order

1. **Verification.** The nearest thing to a service. N-9-26 asks for independent verification organisations and for kill switches whose efficacy is verified on an ongoing basis. The pipeline, with a second, independent auditor, is a working instance of the method such organisations need. It is only credible if the auditor is independent (§6.3).
2. **The corrigibility design rules, tested on the lab's agents.** A research collaboration with a clear kill criterion.
3. **The $\Omega = 1$ normaliser.** A cheap engineering test with a binary answer.
4. **Raising an agent.** A research line that uses the owner's background, and the one least developed in evidence.

In every case the money comes as a salary, a fellowship, a grant or a contract for the tests. There is no licence revenue. The model of §5 keeps those streams, and its middle case assumes no commercial contract at all.

10.5 What the lab gets by engaging now

- **Its tests are designed before they run.** Each of the decisive tests above is written as a pre-registration with a surrogate gate, and every result is recorded, including the failures.
- **Early sight of a method that could remove a sweep** from continual fine-tuning, or turn "kill switch" into measurable quantities, and an honest early stop if it does not.
- **A researcher whose training is in how learners are raised.** That is the angle the lab's character and constitution work needs, and the one this repository's safety results keep pointing to.

10.6 Pitch discipline

say	do not say
"in synthetic worlds, a pause that loses nothing gives an agent on its own clock no reason to resist, at every size we tried"	"we solved shutdown resistance"
"the normaliser survived a 16-fold change of units that broke a tuned weight, on a synthetic problem; on real streams it passed provisionally and failed one replication carrier"	"our continual-learning method beats the state of the art"
"CRR is a synthesis; its own commitments landed on known results in 56 of 80 checkable cases, three of them possibly new"	"CRR is confirmed in fifty fields"
"no result is yet PASS-1; here is the test that would make one"	any number not printed by a committed script

In plain words. If a lab asks "what are you selling?", the honest answer is: a careful way of testing claims, two ideas worth testing on real AI systems, and a short list of experiments that would show quickly whether the ideas work. The careful testing is ready now. The ideas are promising in small pretend worlds. The lab would be paying to find out, fast and honestly, whether they work in the real ones.

11. The continual-learning method and the EPO application, in money terms

Added at the owner's request on 2026-09-23 (prompt-log entry 103). The owner asked for three things:

- to price the continual-learning benchmark results if they hold on GPU and LLM architectures;
- to re-read Continuous_Learning/CONTINUOUS_LEARNING.md for what the pipeline recorded as failures;
- to bring in the owner's European patent application.

The numbers come from a second model, `model/cl_patent.py`, whose output is pinned as `model/cl_patent.txt` (Appendix B). It is separate from the model of \$5, so nothing quoted earlier changes.

11.1 What the continual-learning record actually shows

The EWC arm of the rule was tested on nine unseen carriers in two pre-registered studies. In EQ2 and EQ3 together, the rule at $\Omega = 1$ was not behind the in-sample-tuned weight on 8 of 9 carriers. It was ahead on every one of those 8, by +0.6886 to +4.0625 points.


The one carrier it was behind, fars, is a defect of the pipeline, not a verdict on the rule. The pre-registered subsample left fars with class counts [0, 2086, 746, 15, 991, 688, 430, 45]: one class empty, so the first task had one class. On that carrier every arm of the EWC family sat at chance. Set that carrier aside as uninformative and the rule was not behind on 8 of 8 informative carriers. That reading is **not** a ledger verdict. Under R6 the FAIL stands, and the class floor that would test it is written into EQ4.

The failures the pipeline recorded fall into four kinds; one row, EQX-2, sits in two of them (CONTINUOUS_LEARNING.md §7.5; `cl_patent.txt` [1]):

kind	rows	what it means for money
a pre-registered "every carrier / every cell" criterion missed on one unit	4: EQX-2, EQ3-1, EQ3-C, EQ3-A	mostly the strictness of the test; but EQ3-A is real: at a narrow network (width 64) the rule fell behind
a pre-registered control did not hold	3: EQ2-4, EQ3-3, EQ3-4	the *explanation* was wrong (the rule was predicted to be harmful on constraint-type past terms and often was not); the measurement was not
the theory value $\Omega = 1$ did not appear	3: EQX-2, EQ2-6, EQ3-6	commercially this is the good news: Ω is a plateau (9 of 9 grid points within a step on four carriers), so the dial needs no tuning either
a pipeline defect, or a reduction to a constant	2: EQ2R-VOID, EQX-1	the void cost data, not evidence; the reduction is real: on replay the rule is a constant and worth nothing

One correction to the earlier framing. The owner described the result as "the equanimity finding for optimal learning". The mathematics says otherwise (CONTINUOUS_LEARNING.md §4.3):

- $\Omega = 1$ is not a Bayes optimum. At exact gradients it stops wherever it first meets the Pareto curve.
- What the rule has is a bounded step. It can run the penalty past the point where a fixed weight diverges, and it needs no sweep.
- Its value is therefore *tuning-free and non-divergent*, not *optimal*, and that is what the estimate prices.



The EWC arm on the nine unseen carriers: not behind on 8 of 9 as registered; the one miss is the carrier the pipeline's subsample broke.

The properties that carry value, each from its pinned output:

- **Compute.** On seen data, the rule needed $7.200e+04$ sample passes against the tuned weight's $1.224e+06$ over 17 configurations. That is a ratio of 0.0588: a saving of 0.9412 of that sweep.
- **Invariance.** Across learning rate and batch size the rule was not behind in 5 of 5 cells (EQ3-I, PASS-0), while the tuned weight moved from 300 to 3000.
- **Scale.** On a synthetic problem under SGD, a 16-fold change in the past term's scale left the rule at 0.660 of the re-tuned weight's total loss. The fixed weight at its old setting diverged.
- **The Adam caution, the most important line for LLMs.** Under Adam, the fixed weight is scale-robust too. The rule's scale advantage is an SGD property. Most large models are trained with Adam-family optimisers.
- **Poisoned data.** The bounded rule EQ-B was ahead of the plain rule under poisoned batches on 6 of 6 carriers (by 2.64 to 20.75 points). These were seen carriers, the run was exploratory, and it has no verdict. EQ4 is the held-out test.

In plain words. On nine new datasets the rule did as well as or better than a carefully hand-tuned setting eight times. The ninth time, our own sampling mistake had broken the dataset, and every method failed on it. So the rule works on this kind of problem, it saves the search for the right setting, and it never blows up. The catch is that big AI models are usually trained with a different kind of optimiser, and in our own maths check that optimiser already gives the hand-tuned setting most of the same protection.

11.2 What it would be worth if it held at GPU and LLM scale

The model. The value created is the compute a lab saves by not sweeping one weight. It is computed as:

- the experiment compute of the adopting labs, drawn between \$4,500,000,000 (OpenAI's 2024 experiments, as reported) and \$71,428,571,429 (two labs at a reported \$50B each in 2026, times the 2024 R&D share: an assumption);
- times the share of that compute spent sweeping penalty, KL or anchor weights, an assumption drawn between 0.001 and 0.02;
- times the share of a sweep saved, between 0.8235 (a lab still runs 3 confirming configurations) and 0.9412 (the pinned ratio);
- grown up to 2.2000 a year (the reported training-cost projection).

The technical tree.


- $P(\text{a GPU study reaches PASS-1}) = 0.2500$, by Laplace's rule on the rule's own record of 0 PASS-1 in 2 held-out studies.
- $P(\text{the advantage survives Adam-trained LLMs, given a GPU pass}) = 0.20$. This is an assumption, set low because of the Adam caution.
- $P(\text{labs adopt, given an LLM pass}) = 0.30$, an assumption.

- Together: $P(\text{adopted by 2029}) = 0.0150$.

The results.

quantity	lower	middle	upper
value in the field, 2029, if adopted	\$38,285,915	\$232,739,581	\$1,413,980,686
value in the field, 2030, if adopted	\$48,481,151	\$347,819,511	\$2,484,196,052
value in the field, 2029–2030, if adopted	\$95,290,461	\$580,702,405	\$4,341,466,855


Weighted by the chance of getting there, the mean over all simulated futures for 2029–2030 is \$21,289,825. In 0.0143 of the futures the rule is adopted at all.



If the rule transferred and were adopted, the saved compute is large; weighted by the chance that it does, it is modest.

What moves it. Whether the advantage survives Adam, and how much of their compute labs spend sweeping these weights:

change	mean value created 2029–2030	P(adopted)
registered	\$21,289,825	0.0150
GPU pass at the PASS-0 level (1 of 2)	\$44,956,948	0.0300
LLM transfer 0.05 (Adam removes the advantage)	\$4,030,742	0.0037
LLM transfer 0.5	\$56,277,153	0.0375
sweep share 0.0005 to 0.005	\$6,632,646	0.0150
sweep share 0.005 to 0.05	\$66,326,465	0.0150



What moves the continual-learning value.

How this compares with the estimate in chat of 2026-09-22. That estimate is recorded in CONTINUOUS_LEARNING.md \$9 as a chat figure, not a repository number: "a few million to a few tens

of millions of dollars a year of saved compute under adoption". The committed model differs in two directions:

- **Higher if adopted.** Its conditional-on-adoption values are higher, because 2026 lab compute is reported at about \$50B a lab.
- **Lower once weighted.** Its probability-weighted value is lower than any "under adoption" figure, because the tree now prices the Adam caution and the record's 0 PASS-1.

The fair single sentence is: **worth hundreds of millions of dollars a year to the field if it works at scale, and \$21,289,825 over 2029-2030 in expectation today.**

In plain words. If this trick worked inside the training of the biggest AI models, it could save those companies a lot of computer time, hundreds of millions of dollars a year. But there is a good chance it won't work there, because those models use a different kind of learning step. When you multiply "how much it could be worth" by "how likely it is", you get about twenty million dollars over two years. That is real money for the field, but most of it would be saved by the companies, not paid to us.

11.3 What it could mean for other domains

domain	the same shape?	evidence in this repository	status
RLHF: the KL penalty to a reference policy	yes (β is the weight on the past)	none; the fair test is named: three KL arms (fixed β , an adaptive controller, $\Omega = 1$ on the KL gradient) under SGD and Adam	untested
continual pretraining with an anchor (L2-SP, EWC)	yes	none at scale; the Adam caution applies	untested
safety preservation under fine-tuning (anchor constraints)	yes	the rule amplifies an extreme batch, so it is not a safety mechanism by itself; EQ-B addresses that	exploratory
on-device and edge learning with SGD	yes, and SGD is where the property lives	nearest to the tested instrument	untested; the most natural market
poisoned or adversarial data streams	yes	EQ-B ahead on 6 of 6 seen carriers	exploratory; EQ4 pending
federated personalisation	partly	gated and CLOSED: one global weight is never a step behind	no value found
replay-based continual learning	no (units already match)	EQX: reduces to a constant	no value
adaptive filtering (Kalman)	a fixed gain	$\Omega = 1$ gives a gain that is optimal at one speed only	no value

11.4 The EPO application and the applied use case

What the owner reports (owner's statement):

- a European patent application filed in August 2025, for CRR as a physical systems processor;
- an examiner who has agreed the claims in principle, and has asked for an applied use case.

What the EPO looks for (information, not legal advice; sources in Appendix E). Under the Guidelines for Examination, G-II 3.3, a mathematical method contributes technical character when it is applied to a specific field of technology or adapted to a specific technical implementation. "Controlling a technical system" in general is not enough. The Guidelines' own AI examples include a neural network in a heart-monitoring apparatus that identifies irregular heartbeats, and the classification of images, audio or speech from low-level features. Where a technical effect is established, computational efficiency counts toward inventive step.

Three facts that constrain any use case, stated plainly:

1. **The pending application can only be amended within its content as filed** (Art. 123(2) EPC). A use case must already be disclosed in the August 2025 text, or be directly and unambiguously derivable from it. The repository's results can support a technical effect the application already encompasses: the Enlarged Board's decision G 2/21 allows post-filing evidence for such an effect (named, not fetched). They cannot add a new one.
2. **This repository is public, and has been since 2026-09-15.** In Europe there is no general

grace period, so everything first published here is prior art against any *new* European filing. That includes the Ω rule and EQ-B, the pause-and-resume construction, the alternans and Bass results, and the bounded-memory estimator. In the US, a one-year grace period covers the inventor's own disclosures, so a US filing on matter first published on 2026-09-15 would have to be made no later than 2027-09-15.

3. **The Paris Convention priority year for filings elsewhere that claim the August 2025 date ended in August 2026.**


Candidate applied use cases, mapped to the evidence (for the owner's attorney to check against the application as filed; not claims):

use case	the CRR ingredient	the evidence here	fit with G-II 3.3
cardiac rhythm monitoring or pacing with a bounded memory of settled beats	A6 (regeneration) with P3 weights	ADDS candidate: the alternans threshold moves from 1.0000 to 2.3333 at $q = 0.4$ (synthesis batch 12, row 2)	close to the Guidelines' own heart-monitoring example; Art. 53(c) excludes diagnostic methods practised on the body, not devices
an interruptible controller for an autonomous machine (robot, vehicle): a pause that loses nothing and an objective on the machine's own active steps	H-L5 (natural time), A3 (the cut has no content)	theorem on synthetic worlds from 12 to 768 states ('AI_Safety/checks/scale.txt')	control of a specific physical system, if claimed specifically
drift-robust estimation for sensor streams (bounded mean in place of accumulated counts)	A6	ADDS candidate: better than accumulation by +0.2025 in mean absolute error (synthesis batch 28, row 4)	signal processing of physical measurements
training controller for neural networks on accelerators: a step-bounded penalty weight that removes a sweep	H-EQ (Ω), EQ-B	PASS-0 on tabular data; compute ratio 0.0588 on seen data; the Adam caution	ML training in general is weak; stronger if tied to a specific technical application or hardware implementation
event segmentation of sensor signals at intrinsic-phase half-turns	A3	mixed: H-L5 failed on measles; A3 undecided on projective carriers	a specific signal-processing use would be needed

How the model prices the patent.

- The model multiplies three probabilities, all assumptions: an applied use case exists within the application as filed (0.50); grant given that use case (0.70, reflecting the owner's report); and the granted claims cover what a lab actually uses (0.10). The product is 0.0350.
- On the continual-learning path alone, a patent holder's mean capture over 2029–2030 is \$1,130. Only 0.0004 of the simulated futures produce any income. In those futures it is \$373,639 / \$1,499,953 / \$5,484,681 (lower / middle / upper).
- The cost side is at least EUR 5,250 in EPO renewal fees through year 7, before attorney, validation and national fees.
- **The honest conclusion.** As a way to capture the continual-learning value, the patent is a long shot. That is mainly because the method was developed and published here after the filing date. Its better use is as the vehicle for an applied use case the application already discloses, with the repository's results as supporting evidence. It is also a credential that strengthens the O-1 file (original contributions) and the pitch to the lab.

The patent path, as the model prices it



The patent path as the model prices it, and the calendar that bears on it.

What to do, in order (information, not legal advice):

1. **Engage a European patent attorney now.** Take the application as filed, the examiner's communication, and this table.
2. **Ask the attorney before the next push** whether to pause publishing new methods in this public repository until any separate filings are decided. Every push is a publication.
3. **Decide by 2027-09-15** whether to file in the US on the matter first published here.

In plain words. The patent is like a claim ticket filed in August 2025. You can't add new things to the ticket now; you can only point to what was already written on it, and show proof that it works. Everything we've written in this project since then is public, and in Europe what's public can't be patented later, even by the person who wrote it. So the patent is most useful as a way to protect a use that the original ticket already describes, like a heart-monitoring device or a machine that can be safely paused. A lawyer who knows European patents needs to look at it soon.

11.5 What §11 changes in the plan

- **The earlier totals do not change.** The project funding and researcher compensation of \$5 stand as computed.
- **The continual-learning method adds value to the field, mostly not to its author.** Its expected value to the field is about \$21 million over 2029-2030 in this model. Most of it would be saved by the labs.
- **The author's route to that value** is a role or grant to run the decisive tests (the GPU study with SGD and Adam arms, and the three-arm KL test), and a patent only where the application as filed covers the use.
- **The next steps of §8 gain three items:** run EQ4; propose the GPU and KL tests as a funded project; take the EPO application to an attorney before the next push.

11.6 Addendum, 2026-09-23 (prompt-log entry 104): new evidence on the continual-learning method

This addendum was written after §11 and changes none of the model's inputs. The owner decides whether they should change. The evidence is in Continuous_Learning/ADAM_AND_PRIOR_ART.md and reports/eq4.md. Every number below is pinned there.

- **Against a finely tuned constant, the rule loses where the scales differ.** On a fine grid of 60 weights, the rule's total is 1.186 to 2.073 times the best constant's at the mismatched scales, and 0.997 to 1.050 where they match (adam_checks_3.txt B5). The earlier advantage (0.660 at $c = 16$) was measured against a coarse 10-point grid.
- **The weight's form is published prior art.** Without its smoothing the rule is the VQGAN adaptive weight of 2020 (adam_checks_2.txt B1: 7.7432 against 7.7434), and the registered smoothing costs at every scale tried.
- **Under Adam the rule has nothing to do** (A1-A3). Under heavy-ball momentum it is finite but behind (A8).
- **On real data (EQ4, plain SGD):**

- the bounded rule is not behind the tuned λ on only 4 of 6 unseen carriers (EQ4-1 FAIL); - under poison a fixed weight with the same clip does as well (EQ4-3 FAIL); - the ER-sum control is violated.

What this means for §11.2. The scenario "the rule holds at GPU and LLM scale" now has no synthetic support for an accuracy gain, and under Adam none is possible. What remains is a tuning-free stable weight under SGD-family training. That is a tuning-cost saving which has not been measured, and it is shared with published methods.

The model's value-if-adopted figures (\$232,739,581 for 2029 and \$347,819,511 for 2030, middle case) price a scenario that the evidence now argues against. Read them as a ceiling on a scenario the evidence does not support, not as a forecast.

For the patent (information, not legal advice). The ratio's form was published in 2020, before the August 2025 filing, so accuracy cannot be the technical effect. The applied use case the owner asked for is examined separately, for the cut $\delta(\text{Now})$ as a rupture detector, in Rupture_Detection/RUPTURE_DETECTION.md.

Appendix A — the code

Alexander Plan/model/projections.py

```
"""Funding scenarios for the CRR research programme, 2026 Q4 to 2030 (Alexander Plan; owner request prompt-log entry 101).
```

```
What this is. A deterministic Monte Carlo model of the money that could flow to the programme (grants, fellowships) and to its researcher (employment), under explicit technical scenarios. It is NOT a forecast of what will happen: every input is either (i) an anchor computed from this repository's own record, (ii) a figure reported by a web-search summary on
```

2026-09-23 (listed with its URL in docs/citations/alexander_plan_2026-09-23.md; summaries, not fetched pages), or (iii) an ASSUMPTION, named as such, with a sensitivity sweep. The lower / middle / upper bounds are the 10th, 50th and 90th percentiles of the simulated totals. Every number the dossier quotes is printed here (R1).

The scenario tree (technical milestones; each resolves at the end of a year):

M1 end 2027 a held-out PASS-1 row (or an AI-safety construction shown on a learned model under a declared test)

M2 end 2028 a replication under a fresh prereg (PASS-2), given M1

M3 end 2029 uptake: a lab, a standard or a funder adopts the method, given M2

S0 = no M1; S1 = M1 only; S2 = M1 and M2; S3 = M1, M2 and M3.

Anchors from the record (docs/notes/2026-09-22_epistemic_status.md §1; row ids checked against ledger/LEDGER.md below):

held-out scored rows: PASS-0 EQ2-1b, EQ3-I; FAIL MEAS2-1, EQX-2, EQ3-1 -> 5 rows, 0 PASS-1

replication attempts: 1 (EQ3-1 replicating EQ2-1b), successes 0

p(PASS-1 per study) by Laplace's rule $(k + 1) / (n + 2)$; Jeffreys $(k + 1/2) / (n + 1)$ in the sensitivity table.

uv run python "Alexander Plan/model/projections.py" > "Alexander Plan/model/projections.txt"

"""

from __future__ import annotations

import re

from pathlib import Path

import numpy as np

ROOT = Path(__file__).resolve().parents[2]

SEED = 20260923; N_DRAW = 20000

YEARS = ("2026Q4", "2027", "2028", "2029", "2030"); YEAR_FRAC = (0.25, 1.0, 1.0, 1.0, 1.0)

----- (i) anchors from the repository's record

HELD_OUT = {"EQ2-1b": "PASS-0", "EQ3-I": "PASS-0", "MEAS2-1": "FAIL", "EQX-2": "FAIL", "EQ3-1": "FAIL"}

N_PASS1 = 0; REPL_ATTEMPTS = 1; REPL_SUCCESSES = 0

EQ3_OMEGA_GRID = 9 # the Ω grid of CLAUDE.md §6 used in EQ3 (nine points)

----- (ii) reported figures (search summaries, 2026-09-23)

REPORTED = {

"ltff_accept": (0.193, "Long-Term Future Fund acceptance rate, May 2023 - March 2024"),

"nsf_rate": (0.19, "NSF funding rate, FY2025 (reported fall from an estimated 0.27)"),

"crf_lo": (5000, "Corrigibility Research Fund, typical grant, low (USD)"),

"crf_hi": (35000, "Corrigibility Research Fund, typical grant, high (USD)"),

"crf_total": (200000, "Corrigibility Research Fund, 2026 total (USD)"),

"aisi_lo_gbp": (50000, "UK AISI Alignment Project, grant low (GBP)"),

"aisi_hi_gbp": (1000000, "UK AISI Alignment Project, grant high (GBP)"),

"aisi_total_gbp": (27000000, "UK AISI Alignment Project, first round total (GBP)"),

"aisi_projects": (60, "UK AISI Alignment Project, first round projects"),

"fmf_total": (5000000, "Frontier Model Forum AI Safety Fund, Dec 2025 cohort, more than (USD)"),

"fmf_grantees": (11, "Frontier Model Forum AI Safety Fund, Dec 2025 cohort, grantees"),

"anth_week": (3850, "Anthropic Fellows stipend per week (USD)"),

"mats_week": (1250, "MATS stipend per week (USD)"),

"sal_lo": (97000, "AI safety researcher salary, US range low (USD)"),

"sal_med": (141000, "AI safety researcher salary, US median (USD)"),

"sal_hi": (187000, "AI safety researcher salary, US range high (USD)"),

"frontier_lo": (600000, "frontier-lab senior individual contributor, total compensation low (USD)"),

"frontier_hi": (900000, "frontier-lab senior individual contributor, total compensation high (USD)"),

"mkt_2025": (3.61e9, "AI safety market 2025 (USD; Research and Markets)"),

"mkt_2026": (4.9e9, "AI safety market 2026 (USD; Research and Markets)"),

"mkt_2030": (16.56e9, "AI safety market 2030 forecast (USD; Research and Markets)"),

"mkt_2030_alt": (13.4e9, "AI safety market 2030 forecast (USD; New Market Pitch)"),

"eval_2025": (1.64e9, "AI safety evaluation market 2025 (USD; InsightAce)"),

"uscis_il29": (1055, "USCIS Form I-129 fee for 0 petitions (USD)"),

"uscis_asylum": (600, "USCIS Asylum Program Fee (USD)"),

"uscis_premium": (2965, "USCIS premium processing, Form I-907 (USD)"),

"ol_approval": (0.94, "0-1 approval rate FY2025 (share of adjudicated petitions)"),

"ol_rfe": (0.187, "0-1 share receiving a Request for Evidence, FY2025 to date"),

}

context figures quoted in the dossier but not used by the model (search summaries, 2026-09-23)

CONTEXT = (

("Palisade 2025 tests: share of trials in which Grok 4 resisted shutdown", "nearly 0.90"),

("Palisade 2025 tests: share for codex-mini", "about 0.40 to 0.50"),

("UK AISI Alignment Project: OpenAI's contribution", "GBP 5.6M"),

("UK AISI Alignment Project: AISI's initial funding", "GBP 15M"),

("Frontier Model Forum AI Safety Fund: established with", "over USD 10M"),

("Long-Term Future Fund: acceptance Jan 2022 - Apr 2023, including desk rejections", "0.374"),

("Anthropic Fellows: compute per fellow per month", "about USD 15k"),

("0-1: share approved after a Request for Evidence", "0.709"),

("0-1: regular processing, mid-2026", "up to about 12.5 months"),

("0-1: premium processing", "action within 15 business days"),

("0 petitions completed, FY2025", "nearly 32,000"),

("0-1A: evidentiary criteria to meet (absent a major internationally recognised award)", "at least 3 of 8"),

("California Executive Order N-9-26 (18 Sept 2026): recommendations due", "16 Nov 2026"),

("EU AI Act, general-purpose models: enforcement powers from 2 Aug 2026; fines up to", "EUR 15M or 0.03 of global turnover, whichever is higher"),

)

----- (iii) assumptions (named; swept below)

ASSUME = {

"fx_usd_per_gbp": (1.30, "USD per GBP, for the AISI grant range (not sourced)"),

"anth_weeks": (16, "an Anthropic-Fellows-type fellowship counted as 16 weeks ('four months')"),

"n_studies_2027": (3, "held-out studies scored by end 2027: EQ4 and Tlx2 (pre-registered) and one AI-safety study"),

"p_uptake": (0.30, "P(M3 uptake by end 2029 | M2): no anchor exists"),

"small_apps": (2, "small-grant applications per year"),

"medium_apps": (1, "medium-grant applications per year, from 2027"),

```

    "fellow_apps": (1, "fellowship applications per year"),
}
# success-rate multipliers on the reported base rates, by programme state (a proposal resting on toy results is weaker
# than one resting on a replicated result): ASSUMPTIONS
MULT = {"pre": 0.5, "S0": 0.25, "S1": 1.0, "S2": 1.5, "S3": 2.0}
FELLOW_P = {"pre": 0.05, "S0": 0.03, "S1": 0.10, "S2": 0.20, "S3": 0.30} # ASSUMPTION: no published rate found
EMPLOY_P = {"pre": 0.5, "S0": 0.5, "S1": 0.7, "S2": 0.8, "S3": 0.9} # ASSUMPTION: P(a paid research role that
year, if not already in one) # ASSUMPTION: P(frontier-lab role | S3,
employed)
RETAIN = 0.9 # ASSUMPTION: P(a paid role continues into the
next year)

def check_anchors():
    led = (ROOT / "ledger" / "LEDGER.md").read_text()
    for rid in HELD_OUT:
        assert re.search(rf"^\| {re.escape(rid)} \|", led, re.M), rid
    note = (ROOT / "docs" / "notes" / "2026-09-22_epistemic_status.md").read_text()
    assert "No row is PASS-1" in note and "Two, both PASS-0" in note
    return len(HELD_OUT), sum(v.startswith("PASS") for v in HELD_OUT.values())

def tree(p_study, p_repl, p_up, n_studies):
    p1 = 1 - (1 - p_study) ** n_studies; p2 = p_repl; p3 = p_up
    return dict(p1=p1, p2=p2, p3=p3, S0=1 - p1, S1=p1 * (1 - p2), S2=p1 * p2 * (1 - p3), S3=p1 * p2 * p3)

def simulate(tr, rate_scale=1.0, seed=SEED, n=N_DRAW):
    R = {k: v[0] for k, v in REPORTED.items()}; A = {k: v[0] for k, v in ASSUME.items()}
    rng = np.random.default_rng(seed)
    m1 = rng.random(n) < tr["p1"]; m2 = m1 & (rng.random(n) < tr["p2"]); m3 = m2 & (rng.random(n) < tr["p3"])
    final = np.where(m3, 3, np.where(m2, 2, np.where(m1, 1, 0)))
    proj = np.zeros((n, len(YEARS))); pers = np.zeros((n, len(YEARS))); employed = np.zeros(n, bool)
    base_rate = 0.5 * (R["ltff_accept"] + R["nsf_rate"])
    for j, (y, f) in enumerate(zip(YEARS, YEAR_FRAC)):
        if y in ("2026Q4", "2027"): state = np.full(n, "pre", dtype=object)
        elif y == "2028": state = np.where(m1, "S1", "S0").astype(object)
        elif y == "2029": state = np.where(m2, "S2", np.where(m1, "S1", "S0")).astype(object)
        else: state = np.where(m3, "S3", np.where(m2, "S2", np.where(m1, "S1", "S0"))).astype(object)
        mult = np.array([MULT[s] for s in state]); fel = np.array([FELLOW_P[s] for s in state]); emp = np.array([EMPLOY_P[s] for
s in state])
        for _ in range(1 if y == "2026Q4" else A["small_apps"]):
            win = rng.random(n) < np.minimum(1.0, base_rate * mult * rate_scale)
            proj[:, j] += win * rng.uniform(R["crf_lo"], R["crf_hi"], n)
        if y != "2026Q4":
            for _ in range(A["medium_apps"]):
                win = rng.random(n) < np.minimum(1.0, base_rate * mult * rate_scale)
                proj[:, j] += win * rng.uniform(R["aisi_lo_gbp"], R["aisi_hi_gbp"], n) * A["fx_usd_per_gbp"]
            for _ in range(A["fellow_apps"]):
                win = rng.random(n) < np.minimum(1.0, fel * rate_scale)
                proj[:, j] += win * R["anth_week"] * A["anth_weeks"]
            fresh = rng.random(n) < np.minimum(1.0, emp * rate_scale); kept = rng.random(n) < RETAIN
            employed = fresh if j == 0 else np.where(employed, kept | fresh, fresh)
            salary = rng.uniform(R["sal_lo"], R["sal_hi"], n)
            frontier = (state == "S3") & (rng.random(n) < FRONTIER_P_S3)
            salary = np.where(frontier, rng.uniform(R["frontier_lo"], R["frontier_hi"], n), salary)
            pers[:, j] = employed * salary * f
    return dict(final=final, proj=proj, pers=pers)

def pct(x): return np.percentile(x, [10, 50, 90])

def money(v): return f"${v:,.0f}"

def main():
    n_held, n_pass0 = check_anchors()
    print("Funding scenarios for the CRR programme, 2026 Q4 to 2030 (Alexander Plan/model/projections.py)")
    print(f"seed {SEED}, {N_DRAW} draws; bounds = 10th / 50th / 90th percentiles (lower / middle / upper)")
    print("[1] Anchors from the repository's record (row ids checked in ledger/LEDGER.md)")
    print(f"held-out scored rows {n_held}: " + ", ".join(f"{k} {v}" for k, v in HELD_OUT.items()) + f"; PASS-0 {n_pass0};
PASS-1 {N_PASS1}")
    p_lap = (N_PASS1 + 1) / (n_held + 2); p_jef = (N_PASS1 + 0.5) / (n_held + 1); p0_lap = (n_pass0 + 1) / (n_held + 2)
    p_repl = (REPL_SUCCESSES + 1) / (REPL_ATTEMPTS + 2)
    print(f"p(PASS-1 per study): Laplace {p_lap:.4f}, Jeffreys {p_jef:.4f}; p(PASS-0 per study), Laplace {p0_lap:.4f}")
    print(f"p(replication | a PASS-1): Laplace on {REPL_SUCCESSES} of {REPL_ATTEMPTS} attempts {p_repl:.4f}")
    print("[2] Reported figures (search summaries 2026-09-23; see docs/citations/alexander_plan_2026-09-23.md)")
    for k, (v, d) in REPORTED.items(): print(f" {k:15s} {v:>16.3f} {d}" if v < 10 else f" {k:15s} {v:>16.0f} {d}")
    print(" context (quoted in the dossier, not used by the model):")
    for d, v in CONTEXT: print(f" {d:} {v}")
    R = {k: v[0] for k, v in REPORTED.items()}; A = {k: v[0] for k, v in ASSUME.items()}
    print(" derived: AISI mean first-round grant " + money(R["aisi_total_gbp"]) / R["aisi_projects"]).replace("$", "GBP ")
    + f"; FMF mean grant at least {money(R['fmf_total']) / R['fmf_grantees']}; Anthropic-type fellowship
{money(R['anth_week'] * A['anth_weeks'])}; "
    f"MATS 10 to 12 weeks {money(R['mats_week'] * 10)} to {money(R['mats_week'] * 12)}")
    cagr = (R["mkt_2030"] / R["mkt_2026"]) ** (1 / 4) - 1
    print(f" derived: AI safety market implied growth 2026 -> 2030 {cagr:.4f} per year; path " + ", ".join(f"2026 + t):
${R['mkt_2026'] * (1 + cagr) ** t / 1e9:.2f}B" for t in range(5)))

```

```

    fee = R["uscis_il29"] + R["uscis_asylum"]
    print(f"    derived: 0-1 USCIS fees {money(fee)} base, {money(fee + R['uscis_premium'])} with premium processing (attorney
fees not included)")
    print("[3] Assumptions (named; swept in [6])")
    for k, (v, d) in ASSUME.items(): print(f"    {k:15s} {v:>8g} {d}")
    print("    success-rate multiplier by state " + ", ".join(f"{k} {v:g}" for k, v in MULT.items()) + f" on the base rate {0.5
* (R['ltff_accept'] + R['nsf_rate']):.4f} (mean of LTFF and NSF)")
    print("    fellowship success by state " + ", ".join(f"{k} {v:g}" for k, v in FELLOW_P.items()) + "; employment by state " +
", ".join(f"{k} {v:g}" for k, v in EMPLOY_P.items())
    + f"; frontier-lab role given S3 and employed {FRONTIER_P_S3:g}; a role continues into the next year {RETAIN:g}")
    tr = tree(p_lap, p_repl, A["p_uptake"], A["n_studies_2027"])
    print("[4] Scenario tree (exact)")
    print(f"    p(M1 by end 2027) = 1 - (1 - {p_lap:.4f})^{A['n_studies_2027']} = {tr['p1']:.4f}; p(M2 | M1) = {tr['p2']:.4f};
p(M3 | M2) = {tr['p3']:.4f}")
    for s, name in (("S0", "no held-out result"), ("S1", "a result, not replicated"), ("S2", "replicated, no uptake"), ("S3",
"replicated and taken up")):
        print(f"    {s} {name:26s} p = {tr[s]:.4f}")
    sim = simulate(tr)
    print("[5] Simulated money (USD), all scenarios together")
    print("    project funding (grants and fellowships) per year: lower | middle | upper | mean")
    for j, y in enumerate(YEARS):
        lo, md, hi = pct(sim["proj"][:, j]); print(f"    {y:7s} {money(lo):>12s} | {money(md):>12s} | {money(hi):>12s} |
{money(sim['proj'][:, j].mean()):>12s}")
        cp = sim["proj"].sum(1); cs = sim["pers"].sum(1)
        lo, md, hi = pct(cp); print(f"    cumulative project funding 2026Q4-2030: lower {money(lo)} | middle {money(md)} | upper
{money(hi)} | mean {money(cp.mean())} | P(zero) {np.mean(cp == 0):.4f}")
        print("    researcher compensation per year: lower | middle | upper | mean")
        for j, y in enumerate(YEARS):
            lo, md, hi = pct(sim["pers"][:, j]); print(f"    {y:7s} {money(lo):>12s} | {money(md):>12s} | {money(hi):>12s} |
{money(sim['pers'][:, j].mean()):>12s}")
            lo, md, hi = pct(cs); print(f"    cumulative compensation 2026Q4-2030: lower {money(lo)} | middle {money(md)} | upper
{money(hi)} | mean {money(cs.mean())}")
        print("    by final scenario: share of draws | cumulative project funding lower / middle / upper | cumulative compensation
lower / middle / upper")
        for k in range(4):
            sel = sim["final"] == k; a = pct(cp[sel]); b = pct(cs[sel])
            print(f"    S{k} {sel.mean():.4f} | {money(a[0])} / {money(a[1])} | {money(b[0])} / {money(b[1])} /
{money(b[2])}")
        print("    mean project funding per year by final scenario: " + ", ".join(
f"S{k} " + ", ".join(f"{money(sim['proj'][:, j].mean()) == k, j}.mean())" for j in range(len(YEARS))) for k in range(4)))
    print("[6] Sensitivity (one change at a time): cumulative project funding middle | upper | P(S3)")
    cells = [{"registered", dict(p=p_lap, up=A["p_uptake"], rs=1.0, ns=A["n_studies_2027"])},
("Jeffreys prior", dict(p=p_jef, up=A["p_uptake"], rs=1.0, ns=A["n_studies_2027"])),
("uptake 0.1", dict(p=p_lap, up=0.1, rs=1.0, ns=A["n_studies_2027"])),
("uptake 0.5", dict(p=p_lap, up=0.5, rs=1.0, ns=A["n_studies_2027"])),
("success rates x0.5", dict(p=p_lap, up=A["p_uptake"], rs=0.5, ns=A["n_studies_2027"])),
("success rates x2", dict(p=p_lap, up=A["p_uptake"], rs=2.0, ns=A["n_studies_2027"])),
("1 study in 2027", dict(p=p_lap, up=A["p_uptake"], rs=1.0, ns=1)),
("6 studies in 2027", dict(p=p_lap, up=A["p_uptake"], rs=1.0, ns=6))]
    for name, c in cells:
        t = tree(c["p"], p_repl, c["up"], c["ns"]); s = simulate(t, rate_scale=c["rs"]); x = s["proj"].sum(1)
        print(f"    {name:20s} middle {money(np.percentile(x, 50)):>12s} | upper {money(np.percentile(x, 90)):>12s} | P(S3)
{t['S3']:.4f}")
    print("[7] Value at stake for the field (illustration, not an estimate): share of the 2030 AI safety market, only in S3")
    for share in (0.0001, 0.001, 0.01):
        v = share * R["mkt_2030"]; print(f"    share {share:g}: {money(v)} a year if S3; times P(S3) {tr['S3']:.4f} = {money(v *
tr['S3'])} expected")
    print("[8] Continual learning: what a no-tuning rule would save, if it worked")
    save = 1 - 1 / EQ3_OMEGA_GRID
    print(f"    a {EQ3_OMEGA_GRID}-point weight grid replaced by one run saves {save:.4f} of that sweep's compute")
    for name, p in (("PASS-0 per study (Laplace)", p0_lap), ("PASS-1 per study (Laplace)", p_lap), ("replicated by end 2028
(tree)", tr["p1"] * tr["p2"])):
        print(f"    expected saving per unit of sweep compute, p(works) = {name} {p:.4f}: {save * p:.4f}")
    print("    status from the ledger: EQX the rule at  $\Omega = 1$  reduces to a fixed replay weight; EQ2-1b PASS-0, fragile; EQ3-1
FAIL on 1 of 6; EQ4 (the bounded rule) not yet run")

if __name__ == "__main__":
    main()

```

Alexander Plan/model/cl_patent.py

""The continual-learning method and the EP0 application in money terms, 2026 Q4 to 2030 (Alexander Plan §11; owner request, prompt-log entry 103). A model, not a forecast. Separate from projections.py, so the numbers already quoted from that model do not move.

What it does.

- [1] Re-reads the continual-learning record from the pinned outputs: the per-carrier results of the two held-out studies on the EWC arm (runs/eq2/score.txt, runs/eq3/score.txt), the fars defect (runs/eq3/exclusions.txt class counts), and the ledger rows the pipeline recorded as failures, sorted into the four kinds of Continuous Learning/CONTINUOUS_LEARNING.md §7.5. It prints the strict (as registered) reading and a pipeline-adjusted reading that sets aside the one carrier on which the whole EWC family sat at chance. The adjusted reading is NOT a ledger verdict; it is what a fresh prereg with a class floor (study EQ4 carries one) would test.
- [2] Prints the properties that carry commercial value, each from its pinned output: the tuning-inclusive compute ratio (runs/eq2r_cc/dryrun_on_eq2_seen.txt, seen data), learning-rate x batch invariance (ledger EQ3-I), scale robustness under SGD (theory/checks/omega_vs_methods.txt), the Adam caution (theory/checks/omega_reprocessed.txt), the Omega plateau (ledger EQ3-P) and the bounded rule on poisoned batches (runs/eq4_dev/summary.txt, seen data, exploratory).
- [3] Simulates the value the rule would create at GPU/LLM scale if it transferred, and what an owner of a patent covering it could capture, under a technical tree whose first branch is anchored on the rule's own held-out record.
- [4] Sensitivity, one assumption at a time.

```

[5] The patent calendar that bears on an applied use case (dates only; information, not legal advice).
Inputs are (i) parsed from pinned repository outputs, (ii) figures reported by web-search summaries on 2026-09-23 (URLs in
docs/citations/alexander_plan_2026-09-23.md; summaries, not fetched pages), or (iii) ASSUMPTIONS, named and swept.
uv run python "Alexander Plan/model/cl_patent.py" > "Alexander Plan/model/cl_patent.txt"
"""
from __future__ import annotations

import datetime as dt
import math
import re
from pathlib import Path

import numpy as np

ROOT = Path(__file__).resolve().parents[2]
SEED = 20260923; N_DRAW = 20000

# ----- (ii) reported figures (search summaries, 2026-09-23)
REPORTED = {
    "openai_compute_2024": (7.0e9, "OpenAI cloud compute spend, 2024 (Epoch AI)"),
    "openai_rnd_2024": (5.0e9, "of which R&D compute, 2024 (Epoch AI)"),
    "openai_experiments_2024": (4.5e9, "of which research, experiments and unreleased models, 2024 (Epoch AI)"),
    "openai_final_runs_2024": (0.5e9, "of which final training runs of released models, about, 2024 (Epoch AI)"),
    "lab_compute_2026": (50e9, "compute spend per frontier lab, 2026, OpenAI and Anthropic each, about (press summaries; sources
differ)"),
    "openai_training_2026": (25e9, "OpenAI training cost projection, 2026 (press summary)"),
    "openai_training_2028": (121e9, "OpenAI training cost projection, 2028 (press summary)"),
    "epo_renewals_to_y7_eur": (5250, "EPO renewal fees for one application through year 7, 2026 rates (EUR)"),
}
CONTEXT = (
    ("RL post-training share of DeepSeek-RL's compute", "about 0.05 (Epoch AI summary)"),
    ("a single RL ablation run", "10,000 to 100,000 GPU hours (search summary)"),
    ("hyperscaler capital expenditure, 2026", "USD 660B to 690B (Futurum summary)"),
    ("EPO official fees", "most up about 0.05 from 1 April 2026"),
    ("European SME patents commercially exploited", "up to two thirds (EPO Patent Commercialisation Scoreboard)"),
    ("European universities' and public research organisations' inventions brought to market", "more than a third (EPO
scoreboard)"),
    ("EPO Guidelines G-II 3.3", "a mathematical method contributes technical character by application to a specific field of
technology "
    "and/or adaptation to a specific technical implementation; 'controlling a technical system' in general is not enough"),
    ("EPC Article 54 / 55", "no general grace period: the applicant's own earlier publication is prior art (Art. 55: six-month "
    "exceptions for evident abuse and recognised exhibitions only)"),
    ("EPC Article 123(2)", "an application may not be amended to contain subject matter extending beyond the content as filed"),
    ("35 U.S.C. 102(b)(1)(A)", "a US one-year grace period for the inventor's own disclosures"),
)

# ----- (iii) assumptions (named; swept in [4])
ASSUME = {
    "p_llm_given_gpu": (0.20, "P(the rule's advantage survives AdamW-trained LLM regimes | a GPU PASS-1): the repository's own
check says the "
    "scale advantage is an SGD property"),
    "p_adapt_given_llm": (0.30, "P(labs adopt | an LLM-scale pass)"),
    "s_reg_lo": (0.001, "share of experiment compute spent sweeping penalty / KL / anchor weights, low"),
    "s_reg_hi": (0.02, "share of experiment compute spent sweeping penalty / KL / anchor weights, high"),
    "confirm_configs": (3, "configurations a lab would still run to confirm the rule (of the instrument's 17)"),
    "rnd_share_2026": (5.0 / 7.0, "R&D share of 2026 lab compute, taken from the 2024 OpenAI split"),
    "n_labs_hi": (2, "frontier labs in the high end of the spend range"),
    "p_use_case_ok": (0.5, "P(an applied use case exists in the EPO application as filed and satisfies Art. 123(2))"),
    "p_grant_given_use": (0.7, "P(grant | such a use case; owner's statement: claims agreed in principle)"),
    "p_cover": (0.10, "P(granted claims cover an implementation a lab actually uses)"),
    "roy_lo": (0.005, "share of the value created a patent holder captures, low"),
    "roy_hi": (0.05, "share of the value created a patent holder captures, high"),
}
YEARS_VALUE = (2029, 2030) # GPU study 2027, LLM test 2028, adoption 2029-2030 (timeline assumption)

def parse_carriers(path, row):
    line = next(l for l in path.read_text().splitlines() if l.startswith(row + " "))
    return [(c, float(d), float(s), int(a), int(b)) for c, d, s, a, b in
        re.findall(r"(\w+):([+-]\d+\.\d+) \(\step (\d+\.\d+)\); seeds not behind (\d)/(\d)\)", line)]

def laplace(k, n): return (k + 1) / (n + 2)

def usd(v): return f"${v:,.0f}"

def evidence():
    print("[1] The continual-learning record, re-read (EWC arm, the two held-out studies)")
    eq2 = parse_carriers(ROOT / "runs/eq2/score.txt", "EQ2-1"); eq3 = parse_carriers(ROOT / "runs/eq3/score.txt", "EQ3-1")
    rows = [("EQ2", *r) for r in eq2] + [("EQ3", *r) for r in eq3]
    for st, c, d, s, a, b in rows:
        print(f"    {st} {c:20s} rule - tuned {d:+.4f} (step {s:.2f}; seeds not behind {a}/{b}) -> {'not behind' if d > -s else
'BEHIND'}")
    counts = re.search(r"fars: .*?class counts \[([d, ]+)\]", (ROOT / "runs/eq3/exclusions.txt").read_text())[1]
    empty = [int(x) for x in counts.split(",")]
    print(f"    fars class counts after the pre-registered subsample: [{counts}] -> {empty} empty class; the first task had one
class and every "
        "EWC-family arm sat at chance there (CONTINUOUS_LEARNING.md §7.4)")

```

```

nb = sum(d > -s for _, _, d, s, _, _ in rows); n = len(rows)
print(f"      strict (as registered): not behind on 38 of {n} carriers; study verdicts EQ2-1 PASS (PASS-0, fragile), EQ3-1 FAIL")
print(f"      pipeline-adjusted (fars set aside as uninformative; NOT a ledger verdict): not behind on 38 of {n - 1} informative carriers")
print(f"      per-carrier Laplace: strict {laplace(nb, n):.4f}, adjusted {laplace(nb, n - 1):.4f}; per-study PASS-1 (0 of 2) {laplace(0, 2):.4f}, "
      f"PASS-0-level (1 of 2) {laplace(1, 2):.4f}")
kinds = {"a pre-registered 'every carrier / every cell' criterion missed on one unit": ["EQX-2", "EQ3-1", "EQ3-C", "EQ3-A"],
        "a pre-registered control did not hold (the mechanism statement, not the measurement)": ["EQ2-4", "EQ3-3", "EQ3-4"],
        "the theory value Omega = 1 did not appear (a plateau instead)": ["EQX-2", "EQ2-6", "EQ3-6"],
        "a pipeline defect or a reduction to a constant": ["EQ2R-VOID", "EQX-1"]}
led = (ROOT / "ledger/LEDGER.md").read_text()
print("      what the pipeline recorded as failures, by kind (CONTINUOUS_LEARNING.md §7.5; ids checked in the ledger):")
for k, ids in kinds.items():
    for i in ids: assert re.search(rf"^\{re.escape(i)} \|", led, re.M), i
    print(f"          {len(ids)} {k}: {' '.join(ids)}")
return laplace(0, 2), laplace(1, 2)

def properties():
    print("[2] The properties that carry commercial value (pinned outputs)")
    dr = (ROOT / "runs/eq2r_cc/dryrun_on_eq2_seen.txt").read_text()
    m = re.search(r"tuned lambda ([\d.e+]+) over (\d+) configs, rule ([\d.e+]+) over 1 config, ratio ([\d.]+)", dr)
    ratio = float(m[4]); nconf = int(m[2])
    print(f"      tuning-inclusive compute (seen data, dry run): rule {m[3]} sample passes against the tuned weight's {m[1]} over {nconf} configurations, "
          f"ratio {ratio:.4f} -> saving {1 - ratio:.4f} of that sweep")
    led = (ROOT / "ledger/LEDGER.md").read_text()
    ei = re.search(r"^\{EQ3-I \|([^\}]*\}){4}([^\}]*\)|", led, re.M)[1].strip()
    print(f"      learning-rate x batch invariance (EQ3-I, PASS-0): {ei[:150]}")
    ov = (ROOT / "theory/checks/omega_vs_methods.txt").read_text()
    om = re.search(r"^\omega s+ s+ s+ [\d.]+ s+ [\d.]+ s+ ([\d.]+) \| s+ ([\d.]+) s+ ([\d.]+)", ov, re.M)
    print(f"      scale robustness under SGD (synthetic): the rule's total {om[1]} ({om[2]} of the tuned weight) at F; {om[3]} ({om[4]} of the "
          f"re-tuned weight) at 16F; scale-robust methods {re.search(r'scale-robust methods [^:]*: (.*)', ov)[1]}")
    orp = (ROOT / "theory/checks/omega_reprocessed.txt").read_text()
    print("      the Adam caution (synthetic): " + re.search(r"computed: (the fixed weight is scale-robust under Adam too[^\n]*)", orp)[1])
    ep = re.search(r"^\{EQ3-P \|([^\}]*\}){4}([^\}]*\)|", led, re.M)[1].strip()
    print(f"      Omega plateau (EQ3-P): {ep} grid points within a step of the best, per carrier")
    dev = (ROOT / "runs/eq4_dev/summary.txt").read_text()
    d = [float(x) for x in re.findall(r"POISON.*?EQ-B k=2: [\d.]+ \| (rule - EQ-B ([+ -][\d.]+); clip", dev)]
    print(f"      bounded rule EQ-B (registered kappa 2) under poisoned batches (seen data, exploratory, no verdict): rule - EQ-B on {len(d)} carriers "
          + ", ".join(f"{x:.2f}" for x in d) + f" -> EQ-B ahead on {sum(x < 0 for x in d)} of {len(d)}; study EQ4 (held-out not yet run)")
    return ratio, nconf

def simulate(p_gpu, A, ratio, nconf, rng, p_llm=None, s_lo=None, s_hi=None, p_cover=None, n=N_DRAW):
    R = {k: v[0] for k, v in REPORTED.items()}
    p_llm = A["p_llm_given_gpu"] if p_llm is None else p_llm; s_lo = A["s_reg_lo"] if s_lo is None else s_lo
    s_hi = A["s_reg_hi"] if s_hi is None else s_hi; p_cover = A["p_cover"] if p_cover is None else p_cover
    ok = (rng.random(n) < p_gpu) & (rng.random(n) < p_llm) & (rng.random(n) < A["p_adopt_given_llm"])
    e_lo = R["openai_experiments_2024"]; e_hi = A["n_labs_hi"] * R["lab_compute_2026"] * A["rnd_share_2026"]
    e26 = np.exp(rng.uniform(math.log(e_lo), math.log(e_hi), n))
    g = np.exp(rng.uniform(0.0, 0.5 * math.log(R["openai_training_2028"] / R["openai_training_2026"]), n))
    s = np.exp(rng.uniform(math.log(s_lo), math.log(s_hi), n))
    f = rng.uniform(1 - A["confirm_configs"] / nconf, 1 - ratio, n)
    vals = {y: e26 * g ** (y - 2026) * s * f for y in YEARS_VALUE} # value created in the field if adopted
    created = sum(vals.values()) * ok
    pat = (rng.random(n) < A["p_use_case_ok"]) & (rng.random(n) < A["p_grant_given_use"]) & (rng.random(n) < p_cover)
    roy = np.exp(rng.uniform(math.log(A["roy_lo"]), math.log(A["roy_hi"]), n))
    owner = created * pat * roy
    return dict(ok=ok, vals=vals, created=created, owner=owner, pat=pat)

def pct(x): return np.percentile(x, [10, 50, 90])

def main():
    print("The continual-learning method and the EPO application in money terms (Alexander Plan/model/cl_patent.py)")
    print(f"seed {SEED}, {N_DRAW} draws; lower / middle / upper = 10th / 50th / 90th percentiles")
    p_pass1, p_pass0 = evidence(); ratio, nconf = properties()
    print("[3a] Inputs")
    for k, (v, d) in REPORTED.items(): print(f"      reported {k:24s} {v:>18,.0f} {d}")
    for d, v in CONTEXT: print(f"      context {d}: {v}")
    A = {k: v[0] for k, v in ASSUME.items()}
    for k, (v, d) in ASSUME.items(): print(f"      ASSUMED {k:18s} {v:>8.4g} {d}")
    e_hi = A["n_labs_hi"] * REPORTED["lab_compute_2026"][0] * A["rnd_share_2026"]
    print(f"      derived: experiment compute of adopting labs in 2026, log-uniform from {usd(REPORTED['openai_experiments_2024'][0])} to {usd(e_hi)}; "
          f"yearly growth up to {math.sqrt(REPORTED['openai_training_2028'][0] / REPORTED['openai_training_2026'][0]):.4f}; share of a sweep saved "
          f"{1 - A['confirm_configs'] / nconf:.4f} to {1 - ratio:.4f}")
    print(f"      derived: P(GPU PASS-1) = Laplace on the rule's held-out studies, 0 PASS-1 of 2 = {p_pass1:.4f}")
    p_tech = p_pass1 * A["p_llm_given_gpu"] * A["p_adopt_given_llm"]
    p_pat = A["p_use_case_ok"] * A["p_grant_given_use"] * A["p_cover"]

```

```

    print(f"    derived: P(adopted by 2029) = {p_pass1:.4f} x {A['p_llm_given_gpu']:.2f} x {A['p_adopt_given_llm']:.2f} =
{p_tech:.4f}); "
        f"P(a granted patent covers the used implementation) = {A['p_use_case_ok']:.2f} x {A['p_grant_given_use']:.2f} x
{A['p_cover']:.2f} = {p_pat:.4f}")
    rng = np.random.default_rng(SEED); sim = simulate(p_pass1, A, ratio, nconf, rng)
    ok = sim["ok"]
    print("[3b] Value the rule would create in the field (saved sweep compute), USD")
    for y in YEARS.VALUE:
        lo, md, hi = pct(sim["vals"][y]); print(f"        {y}, if adopted:                lower {usd(lo)} | middle {usd(md)} | upper
{usd(hi)} | mean {usd(sim['vals'][y].mean())}")
        lo, md, hi = pct(sim["created"])
        print(f"        2029-2030 cumulative, all futures: lower {usd(lo)} | middle {usd(md)} | upper {usd(hi)} | mean
{usd(sim['created'].mean())} | "
            f"share of futures with any value {ok.mean():.4f}")
        lo, md, hi = pct(sim["created"][ok]); print(f"        2029-2030 cumulative, adopted futures only: lower {usd(lo)} | middle
{usd(md)} | upper {usd(hi)}")
        print("[3c] What a patent holder could capture, USD (2029-2030)")
        lo, md, hi = pct(sim["owner"])
        print(f"    all futures: lower {usd(lo)} | middle {usd(md)} | upper {usd(hi)} | mean {usd(sim['owner'].mean())} | share of
futures with any income "
            f"{(sim['owner'] > 0).mean():.4f}")
        sel = sim["owner"] > 0
        if sel.any():
            lo, md, hi = pct(sim["owner"][sel]); print(f"        futures with income only: lower {usd(lo)} | middle {usd(md)} | upper
{usd(hi)}")
        print(f"    cost side: EPO renewal fees through year 7 for one application EUR {REPORTED['epo_renewals_to_y7_eur'][0]:.0f}
(attorney, validation and "
            "national fees not included)")
        print("[4] Sensitivity (one change at a time): mean value created 2029-2030 | mean patent-holder capture | P(adopted)")
        cells = [{"registered", {}}, {"GPU pass at the PASS-0 level (1 of 2)", dict(p_gpu=p_pass0)},
            ("LLM transfer 0.05 (Adam removes the advantage)", dict(p_llm=0.05)), ("LLM transfer 0.5", dict(p_llm=0.5)),
            ("sweep share 0.0005 to 0.005", dict(s_lo=0.0005, s_hi=0.005)), ("sweep share 0.005 to 0.05", dict(s_lo=0.005,
s_hi=0.05)),
            ("patent covers 0.02", dict(p_cover=0.02)), ("patent covers 0.3", dict(p_cover=0.3))]
        for name, kw in cells:
            pg = kw.pop("p_gpu", p_pass1); r = simulate(pg, A, ratio, nconf, np.random.default_rng(SEED), **kw)
            pt = pg * kw.get("p_llm", A["p_llm_given_gpu"]) * A["p_adopt_given_llm"]
            print(f"    {name:46s} {usd(r['created'].mean()):>16s} | {usd(r['owner'].mean()):>12s} | {pt:.4f}")
        print("[5] The patent calendar that bears on an applied use case (dates; information, not legal advice)")
        filed = dt.date(2025, 8, 1); public = dt.date(2026, 9, 15)
        print(f"        EPO application filed August 2025 (owner's statement; day not given, the 1st used for arithmetic):
{filed.isoformat()[7:]}")
        print(f"    Paris Convention priority year for filings elsewhere claiming it: ended {dt.date(2026, 8, 1).isoformat()[7:]}")
        print(f"    repository public since {public.isoformat()} (GitHub API: created_at 2026-09-15T04:06:50Z, visibility public)")
        print(f"    EPO: anything first published in the repository is prior art for any NEW European filing (no grace period); the
pending "
            "application can only be amended within its content as filed (Art. 123(2))")
        print(f"    US: a filing on repository-disclosed matter would have to be made within one year of its first disclosure, i.e.
no later than "
            f"{dt.date(2027, 9, 15).isoformat()} for matter first published on {public.isoformat()}")
        print(f"    a European patent's term runs 20 years from filing: to {dt.date(2045, 8, 1).isoformat()[7:]}")

if __name__ == "__main__":
    main()

```

Alexander Plan/build/make_figures.py

```

"""Figures for Alexander Plan/ALEXANDER_PLAN.md (owner request prompt-log entry 101). Every data figure is drawn from a pinned
output (Alexander Plan/model/projections.txt; AI_Safety/checks/scale.txt), parsed by the regular expressions below; nothing
is recomputed except unit conversions, which are printed. The two planning figures (A08 the 0-1 timeline, A09 the evidence
map) are drawings of a plan and of a first reading, not data, and say so on their faces. Every number placed on a figure is
printed to stdout, pinned as Alexander Plan/figures/figures.txt and CI-checked (R1). Palette: the validated reference
instance used for AI_Safety/ (categorical slots in fixed order; identity never rests on colour alone).
Run: uv run python "Alexander Plan/build/make_figures.py" > "Alexander Plan/figures/figures.txt"
"""

from __future__ import annotations

import re
from pathlib import Path

import matplotlib
matplotlib.use("Agg")
import matplotlib.pyplot as plt # noqa: E402
import numpy as np # noqa: E402
from matplotlib.patches import FancyBboxPatch # noqa: E402

ROOT = Path(__file__).resolve().parents[2]
DOC = ROOT / "Alexander Plan"; OUT = DOC / "figures"
SURFACE, INK, INK2, GRID, MUTED = "#fcfcfb", "#0b0b0b", "#52514e", "#e6e5e1", "#8a8983"
SC = {"S0": MUTED, "S1": "#2a78d6", "S2": "#eda100", "S3": "#1baf7a"}
SNAME = {"S0": "S0 no held-out result", "S1": "S1 a result, not replicated", "S2": "S2 replicated, no uptake", "S3": "S3
replicated and taken up"}
plt.rcParams.update({"figure.facecolor": SURFACE, "axes.facecolor": SURFACE, "axes.edgecolor": INK2, "axes.labelcolor": INK,
    "xtick.color": INK2, "ytick.color": INK2, "text.color": INK, "axes.grid": True, "grid.color": GRID,
    "grid.linewidth": 0.6, "axes.spines.top": False, "axes.spines.right": False, "font.size": 9,
    "axes.titlesize": 10, "legend.frameon": False, "lines.linewidth": 2, "savefig.dpi": 200, "axes.axisbelow":
True})
P = (DOC / "model" / "projections.txt").read_text(); SCALE = (ROOT / "AI_Safety" / "checks" / "scale.txt").read_text()
YEARS = ("2026Q4", "2027", "2028", "2029", "2030")

```

```

def usd(s): return float(s.replace("$", "").replace(",", ""))

def save(fig, name):
    fig.savefig(OUT / name, bbox_inches="tight"); plt.close(fig); print(f"[figure] {name}")

def k(v):
    if v >= 1e6: return f"${v / 1e6:,.2f}M"
    return f"${v / 1000:,.1f}k" if (v < 1e5 and v % 1000) else f"${v / 1000:,.0f}k"

# ----- A01 the scenario tree
def a01_tree():
    m = re.search(r"p\((M1 by end 2027\) = .*? = ([\d.]+); p\((M2 \setminus M1\) = ([\d.]+); p\((M3 \setminus M2\) = ([\d.]+)", P)
    p1, p2, p3 = (float(x) for x in m.groups())
    ps = {s: float(re.search(rf" {s} .*?p = ([\d.]+)", P)[1]) for s in SC}
    fig, ax = plt.subplots(figsize=(10.2, 4.4)); ax.axis("off"); ax.set_xlim(0, 10.2); ax.set_ylim(0, 4.4)
    def box(x, y, t, c=INK2, w=1.9, h=0.62):
        ax.add_patch(FancyBboxPatch((x - w / 2, y - h / 2), w, h, boxstyle="round,pad=0.02,rounding_size=0.08", fc=SURFACE,
ec=c, lw=1.6))
        ax.text(x, y, t, ha="center", va="center", fontsize=7.8)
    def arrow(x0, y0, x1, y1, t, dy=0.0):
        ax.annotate("", (x1, y1), (x0, y0), arrowprops=dict(arrowstyle="-|>", color=INK2, lw=1))
        ax.text((x0 + x1) / 2, (y0 + y1) / 2 + dy, t, ha="center", va="center", fontsize=7.4, color=INK2, bbox=dict(fc=SURFACE,
ec="none", pad=0.6))
    box(1.1, 2.9, "now: toy results,\nno held-out PASS-1")
    box(3.6, 3.6, "M1 end 2027:\nna held-out result"); box(3.6, 1.2, SNAME["S0"].replace("S0 ", "S0\n"), SC["S0"])
    arrow(2.05, 3.05, 2.65, 3.5, f"{p1:.4f}"); arrow(2.05, 2.75, 2.65, 1.35, f"{1 - p1:.4f}")
    box(6.1, 3.9, "M2 end 2028:\nrepeated (PASS-2)"); box(6.1, 2.5, SNAME["S1"].replace("S1 ", "S1\n"), SC["S1"])
    arrow(4.55, 3.7, 5.15, 3.85, f"{p2:.4f}", dy=0.22); arrow(4.55, 3.45, 5.15, 2.65, f"{1 - p2:.4f}")
    box(8.8, 3.9, SNAME["S3"].replace("S3 ", "S3 (M3)\n"), SC["S3"]); box(8.8, 2.7, SNAME["S2"].replace("S2 ", "S2\n"),
SC["S2"])
    arrow(7.05, 3.95, 7.85, 3.95, f"{p3:.4f}", dy=0.2); arrow(7.05, 3.7, 7.85, 2.85, f"{1 - p3:.4f}")
    for s, (x, y) in {"S0": (3.6, 0.6), "S1": (6.1, 1.9), "S2": (8.8, 2.1), "S3": (8.8, 4.35)}.items():
        ax.text(x, y, f"{s}: p = {ps[s]:.4f}", ha="center", fontsize=8.4, color=SC[s] if s != "S0" else INK2, fontweight="bold")
    ax.text(0.05, 0.1, "p(M1) and p(M2) are Laplace estimates from this repository's own record (5 held-out rows, 0 PASS-1; 1
replication, 0 successes); p(M3) is an assumption.",
        fontsize=7.2, color=INK2)
    ax.set_title("Where the programme could be by 2030: four technical scenarios", loc="left")
    print(f"[A01] p1 {p1:.4f} p2 {p2:.4f} p3 {p3:.4f}; " + " ".join(f"{s} {v:.4f}" for s, v in ps.items()))
    save(fig, "A01_scenario_tree.png")

# ----- A02 project funding per year
def a02_yearly():
    line = re.search(r"mean project funding per year by final scenario: (.*)\n", P)[1]
    per = {s: [usd(x) for x in re.search(rf"{s} ((?:\d[\d,]+ ?){5})", line)[1].split()] for s in SC}
    rows = {y: re.search(rf" {re.escape(y)}\s+(\d[\d,]+) \setminus\s+(\d[\d,]+) \setminus\s+(\d[\d,]+) \setminus\s+(\d[\d,]+)", P) for y in YEARS}
    up = [usd(rows[y][3]) for y in YEARS]; mean = [usd(rows[y][4]) for y in YEARS]
    fig, ax = plt.subplots(figsize=(7.4, 3.8)); x = np.arange(len(YEARS))
    ax.fill_between(x, 0, up, color=GRID, alpha=0.9, label="all scenarios: up to the upper bound (90th percentile)")
    ax.plot(x, mean, color=INK, lw=1.4, ls="--", label="all scenarios: mean")
    for s in SC:
        ax.plot(x, per[s], color=SC[s], marker="o", markersize=3.5, label=SNAME[s] + " (mean)")
    ax.set_xticks(x); ax.set_xticklabels(YEARS); ax.set_ylabel("project funding per year (USD)")
    ax.yaxis.set_major_formatter(matplotlib.ticker.FuncFormatter(lambda v, _ : k(v)))
    ax.legend(fontsize=7, loc="upper left"); ax.set_title("Grants and fellowships per year: the middle is zero, the mean is
carried by a few large awards", loc="left", fontsize=9.2)
    print("[A02] upper " + " ".join(f"{v:,.0f}" for v in up) + "; mean " + " ".join(f"{v:,.0f}" for v in mean))
    for s in SC: print(f"[A02] {s} mean " + " ".join(f"{v:,.0f}" for v in per[s]))
    save(fig, "A02_project_funding_per_year.png")

# ----- A03 cumulative, by scenario
def a03_cumulative():
    rows = {s: re.search(rf" {s} ([\d.]+) \setminus ([\d.]+) / ([\d.]+) \setminus ([\d.]+) \setminus ([\d.]+) / ([\d.]+) / ([\d.]+)",
P) for s in SC}
    fig, axes = plt.subplots(1, 2, figsize=(10.4, 3.7))
    for ax, off, title in ((axes[0], 2, "project funding, 2026 Q4 to 2030"), (axes[1], 5, "researcher compensation, 2026 Q4 to
2030")):
        for i, s in enumerate(SC):
            lo, md, hi = (usd(rows[s][off + j]) for j in range(3))
            ax.plot([i, i], [lo, hi], color=SC[s], lw=7, alpha=0.45, solid_capstyle="butt"); ax.plot(i, md, "o", color=SC[s],
markersize=8, mec=INK, mew=0.6)
            ax.text(i + 0.12, md, k(md), fontsize=7.4, va="center"); ax.text(i, hi, k(hi), fontsize=7, ha="center", va="bottom",
color=INK2)
            print(f"[A03] {title} {s}: lower {lo:,.0f} middle {md:,.0f} upper {hi:,.0f}")
        ax.set_xticks(range(4)); ax.set_xticklabels([SNAME[s].replace(" ", "\n", 1) for s in SC], fontsize=7.2)
        ax.yaxis.set_major_formatter(matplotlib.ticker.FuncFormatter(lambda v, _ : k(v))); ax.grid(axis="x", visible=False)
        ax.set_title(title + ": lower to upper, dot = middle", loc="left", fontsize=8.8)
        shares = " ".join(f"{s} {float(rows[s][1]):.4f}" for s in SC); print(f"[A03] shares of draws {shares}")
    fig.tight_layout(); save(fig, "A03_cumulative_by_scenario.png")

# ----- A04 sensitivity
def a04_sensitivity():
    rows = re.findall(r" ([\d.]+\n)*\s+middle\s+(\d[\d,]+) \setminus upper\s+(\d[\d,]+) \setminus P(S3\ ([\d.]+)", P)
    fig, ax = plt.subplots(figsize=(7.4, 3.6)); y = np.arange(len(rows))[:-1]

```

```

for yi, (name, md, hi, s3) in zip(y, rows):
    md, hi = usd(md), usd(hi); c = INK if name == "registered" else "#2a78d6"
    ax.plot([md, hi], [yi, yi], color=c, lw=1.2, alpha=0.6); ax.plot(md, yi, "o", color=c, markersize=6); ax.plot(hi, yi,
"|", color=c, markersize=12, mew=2)
    ax.text(hi * 1.03, yi, f"P(S3) {float(s3):.4f}", va="center", fontsize=7.2, color=INK2)
    print(f"[A04] {name}: middle {md:,.0f} upper {hi:,.0f} P(S3) {float(s3):.4f}")
ax.set_yticks(y); ax.set_yticklabels([r[0] for r in rows], fontsize=7.6); ax.set_xscale("log")
ax.xaxis.set_major_formatter(matplotlib.ticker.FuncFormatter(lambda v, _ : k(v)))
ax.set_xlabel("cumulative project funding 2026 Q4 to 2030 (dot = middle, bar = upper), log scale"); ax.set_xlim(5e3, 4e6)
ax.set_title("What moves the estimate: success rates most, the uptake guess least", loc="left")
save(fig, "A04_sensitivity.png")

# ----- A05 the market
def a05_market():
    path = re.findall(r"(\d{4}): \$(\d{.})+B", re.search(r"implied growth .*?\n", P)[0]); g = float(re.search(r"implied growth
2026 -> 2030 (\d{.})+", P)[1])
    m25 = usd(re.search(r"mkt_2025\s+(\d{.})+", P)[1]); alt = usd(re.search(r"mkt_2030_alt\s+(\d{.})+", P)[1]); ev =
usd(re.search(r"eval_2025\s+(\d{.})+", P)[1])
    fig, ax = plt.subplots(figsize=(7.0, 3.6))
    xs = [int(a) for a, _ in path]; ys = [float(b) for _, b in path]
    ax.plot(xs, ys, color="#2a78d6", marker="o", markersize=4, label=f"AI safety market (Research and Markets), implied growth
{g:.4f} a year")
    ax.plot([2025, 2026], [m25 / 1e9, ys[0]], color="#2a78d6", ls=":", marker="o", markersize=4)
    ax.plot(2030, alt / 1e9, "D", color="#eb6834", label="alternative 2030 forecast (New Market Pitch)")
    ax.plot(2025, ev / 1e9, "s", color="#1baf7a", label="AI safety evaluation market 2025 (InsightAce)")
    for x, y in zip(xs, ys): ax.text(x, y + 0.5, f"${y:.2f}B", ha="center", fontsize=7.2)
    ax.text(2025, m25 / 1e9 + 0.5, f"${m25 / 1e9:.2f}B", ha="center", fontsize=7.2); ax.text(2030.1, alt / 1e9, f"${alt /
1e9:.2f}B", va="center", fontsize=7.2)
    ax.text(2025.1, ev / 1e9 - 0.9, f"${ev / 1e9:.2f}B", fontsize=7.2)
    ax.set_ylabel("USD billions"); ax.set_xlim(2024.6, 2030.7); ax.legend(fontsize=7, loc="upper left")
    ax.set_title("The market, as reported by market-research firms (definitions differ)", loc="left")
    print(f"[A05] path {ys}; 2025 {m25:,.0f}; alt 2030 {alt:,.0f}; eval 2025 {ev:,.0f}; growth {g:.4f}; labels ${m25 /
1e9:.2f}B, ${alt / 1e9:.2f}B, ${ev / 1e9:.2f}B")
    save(fig, "A05_market.png")

# ----- A06 the funding programmes
def a06_programmes():
    fx = float(re.search(r"fx_usd_per_gbp\s+(\d{.})+", P)[1])
    g = lambda key: usd(re.search(rf"{key}\s+(\d{.})+", P)[1])
    fell = usd(re.search(r"Anthropic-type fellowship (\d{.})+", P)[1]); mats = re.search(r"MATS 10 to 12 weeks (\d{.})+ to
(\d{.})+", P)
    fmf = usd(re.search(r"FMF mean grant at least (\d{.})+", P)[1])
    items = [("Corrigibility Research Fund grant", g("crf_lo"), g("crf_hi")), ("MATS stipend, 10 to 12 weeks", usd(mats[1]),
usd(mats[2])),
            ("Anthropic-Fellows-type stipend, 16 weeks", fell, fell), ("FMF AI Safety Fund, mean grant (at least)", fmf, fmf),
            ("UK AISI Alignment Project grant (at USD 1.30 per GBP)", g("aisi_lo_gbp") * fx, g("aisi_hi_gbp") * fx),
            ("AI safety researcher salary, US, a year", g("sal_lo"), g("sal_hi")), ("frontier-lab senior IC, total a year",
g("frontier_lo"), g("frontier_hi"))]
    fig, ax = plt.subplots(figsize=(7.6, 3.6)); y = np.arange(len(items))[:-1]
    for yi, (name, lo, hi), c in zip(y, items, ["#2a78d6", "#2a78d6", "#2a78d6", "#eda100", "#eda100", "#1baf7a", "#1baf7a"]):
        if hi > lo: ax.plot([lo, hi], [yi, yi], color=c, lw=6, alpha=0.7, solid_capstyle="butt")
        else: ax.plot(lo, yi, "o", color=c, markersize=7)
        ax.text(hi * 1.12, yi, (k(lo) + (" if hi == lo else f" to {k(hi)}")).replace("$", r"\$"), va="center", fontsize=7.2)
        print(f"[A06] {name}: {lo:,.0f} to {hi:,.0f}")
    ax.set_yticks(y); ax.set_yticklabels([i[0] for i in items], fontsize=7.4); ax.set_xscale("log"); ax.set_xlim(3e3, 8e6)
    ax.xaxis.set_major_formatter(matplotlib.ticker.FuncFormatter(lambda v, _ : k(v)))
    ax.set_title("Sizes of the money on offer: small grants, fellowships, project grants, salaries", loc="left", fontsize=9.4)
    save(fig, "A06_programmes.png")

# ----- A07 continual learning: the value of a no-tuning rule
def a07_cl():
    save_ = float(re.search(r"saves (\d{.})+ of that sweep", P)[1])
    anchors = re.findall(r"p(works\)) = (.*) (\d{.})+ (\d{.})+\n", P)
    fig, ax = plt.subplots(figsize=(6.8, 3.4)); p = np.linspace(0, 1, 101)
    ax.plot(p, save_ * p, color="#2a78d6", label=f"expected saving = {save_: .4f} x p(works)")
    for (name, pw, v), c in zip(anchors, ("eda100", "#eb6834", "#1baf7a")):
        ax.plot(float(pw), float(v), "o", color=c, markersize=7, label=f"{name}: {float(v):.4f}")
        print(f"[A07] {name}: p {float(pw):.4f} saving {float(v):.4f}")
    ax.set_xlabel("p(the rule works as registered)"); ax.set_ylabel("share of a 9-point sweep's compute saved")
    ax.legend(fontsize=7, loc="upper left"); ax.set_title("If a no-tuning rule worked: the saving is on the weight sweep only",
loc="left")
    print(f"[A07] full saving {save_: .4f}")
    save(fig, "A07_cl_value.png")

# ----- A08 the 0-1 plan (a drawing of a plan)
def a08_timeline():
    steps = [("engage an immigration attorney; confirm the petitioner", 0, 1), ("collect evidence against the criteria", 0.5,
3.5),
            ("expert letters (Levin agreed, per the owner; others asked)", 1, 3.5), ("advisory opinion from a peer group", 2.5,
3.5),
            ("petitioner files I-129 (premium: 15 business days)", 3.5, 4.3), ("respond to any Request for Evidence", 4.3,
5.5),
            ("visa stamping at a US consulate; start", 5, 6.5), ("in parallel: EQ4, T1x2 and a learned-model safety study", 0,
6.5)]
    fig, ax = plt.subplots(figsize=(8.4, 3.4)); y = np.arange(len(steps))[:-1]
    for yi, (name, a, b) in zip(y, steps):

```

```

        c = "#1baf7a" if name.startswith("in parallel") else "#2a78d6"
        ax.barh(yi, b - a, left=a, color=c, alpha=0.75, height=0.55)
        ax.set_yticks(y); ax.set_yticklabels([s[0] for s in steps], fontsize=7.4); ax.set_xlabel("months from engaging an attorney
(a plan, not a forecast)")
        ax.set_xlim(0, 7); ax.grid(axis="y", visible=False)
        ax.set_title("A possible 0-1 sequence: evidence first, then filing", loc="left")
        print("[A08] planning drawing (no data)")
        save(fig, "A08_01_timeline.png")

# ----- A09 the 0-1 evidence map (a first reading)
def a09_criteria():
    crit = [{"1 nationally or internationally recognised awards", "unknown"}, {"2 membership requiring outstanding achievement",
"unknown"},
            {"3 published material about you", "unknown"}, {"4 judging the work of others", "possible"},
            {"5 original contributions of major significance", "possible"}, {"6 scholarly articles", "possible"},
            {"7 critical role at a distinguished organisation", "possible"}, {"8 high salary", "weak"}]
    col = {"possible": "#1baf7a", "unknown": "#eda100", "weak": "MUTED"}
    lab = {"possible": "plausible from what you have told me: document it", "unknown": "not known to me: check", "weak":
"unlikely to carry weight now"}
    fig, ax = plt.subplots(figsize=(8.0, 3.4)); ax.axis("off"); ax.set_xlim(0, 10); ax.set_ylim(0, len(crit) + 1)
    for i, (name, st) in enumerate(crit):
        yy = len(crit) - i
        ax.add_patch(FancyBboxPatch((0.1, yy - 0.38), 0.5, 0.76, boxstyle="round,pad=0.02", fc=col[st], ec=SURFACE))
        ax.text(0.8, yy, name, va="center", fontsize=8); ax.text(6.0, yy, lab[st], va="center", fontsize=7.6, color=INK2)
        ax.text(0.1, 0.25, "A first reading from the owner's statements; the attorney decides. 0-1A needs at least 3 of the 8, each
with documents.", fontsize=7.4, color=INK2)
    ax.set_title("The eight 0-1A evidentiary criteria: where the evidence may come from", loc="left")
    print("[A09] a first reading, not data: " + ", ".join(f"{n.split()[0]} {s}" for n, s in crit))
    save(fig, "A09_01_criteria.png")

# ----- A10-A13 continual learning and the EPO application
(cl_patent.txt)
CLP = (DOC / "model" / "cl_patent.txt").read_text()

def a10_cl_evidence():
    rows = re.findall(r"^\s+(EQ[23]) (\w+)\s+rule - tuned ([+|-]|\d\.)+ \\\(step ([\d\.]+) seeds not behind ([\d\.]+) \\\) -> (not
behind|BEHIND)", CLP, re.M)
    fig, ax = plt.subplots(figsize=(8.2, 3.9)); y = np.arange(len(rows))[:-1]
    for yi, (st, c, d, stp, a, b, lab) in zip(y, rows):
        d, stp = float(d), float(stp); col = "#eb6834" if lab == "BEHIND" else "#2a78d6"
        ax.barh(yi, d, color=col, height=0.6); ax.plot([-stp, -stp], [yi - 0.35, yi + 0.35], color=INK2, lw=1)
        ax.text(d + (0.15 if d >= 0 else -0.15), yi, f"{d:.2f} ({a}/{b} seeds)", va="center", ha="left" if d >= 0 else "right",
        fontsize=7.2)
        print(f"[A10] {st} {c}: {d:.4f} step {stp:.2f} seeds {a}/{b} {lab}")
    ax.axvline(0, color=INK2, lw=0.8); ax.set_yticks(y); ax.set_yticklabels([f"{r[0]} {r[1]}" for r in rows], fontsize=7.8)
    ax.set_xlim(-8, 7); ax.set_xlabel("rule at  $\Omega = 1$  minus the in-sample-tuned weight (accuracy points); tick = minus one
resolvable step")
    fy = y[[r[1] for r in rows].index("fars")]
    ax.text(-7.9, fy - 0.42, "fars: the pre-registered subsample left a class empty; never EWC-family arm sat at chance",
    fontsize=6.8, color="#b3461c", va="top")
    s1 = re.search(r"strict \\\(as registered\\\): not behind on ([\d+]) of ([\d+])", CLP); s2 = re.search(r"not behind on ([\d+]) of
([\d+]) informative", CLP)
    ax.set_title(f"The EWC arm on nine unseen carriers: not behind on {s1[1]} of {s1[2]} (strict), {s2[1]} of {s2[2]}
informative", loc="left", fontsize=9.4)
    ax.grid(axis="y", visible=False); save(fig, "A10_cl_evidence.png")

def a11_cl_value():
    g = lambda lab: re.search(rf"\{lab\}\s*lower ([\d\.]+) \\\ middle ([\d\.]+) \\\ upper ([\d\.]+) ", CLP)
    rows = [{"2029, if adopted", g(r"2029, if adopted:")}, {"2030, if adopted", g(r"2030, if adopted:")},
            {"2029-2030, if adopted", g(r"2029-2030 cumulative, adopted futures only:")}]]
    mean_all = usd(re.search(r"2029-2030 cumulative, all futures: .*?mean ([\d\.]+) ", CLP)[1]); p_ad = float(re.search(r"share
of futures with any value ([\d\.]+) ", CLP)[1])
    fig, ax = plt.subplots(figsize=(8.2, 3.2)); y = np.arange(len(rows))[:-1]
    for yi, (name, m) in zip(y, rows):
        lo, md, hi = usd(m[1]), usd(m[2]), usd(m[3])
        ax.plot([lo, hi], [yi, yi], color="#2a78d6", lw=6, alpha=0.45, solid_capstyle="butt"); ax.plot(md, yi, "o",
        color="#2a78d6", markersize=8, mec=INK, mew=0.6)
        ax.text(hi * 1.15, yi, f"{k(lo)} / {k(md)} / {k(hi)}".replace("$", r"\$"), va="center", fontsize=7.4)
        print(f"[A11] {name}: {lo:.0f} / {md:.0f} / {hi:.0f}")
    ax.axvline(mean_all, color="#eb6834", ls="--", lw=1.2)
    ax.set_ylim(y[-1] - 0.5, y[0] + 0.9)
    ax.text(mean_all * 1.08, y[0] + 0.55, f"probability-weighted mean over all futures (2029-2030): {k(mean_all)}; simulated
share adopted {p_ad:.4f}".replace("$", r"\$"),
    fontsize=7.2, color="#b3461c")
    ax.set_xscale("log"); ax.set_xlim(8e6, 4e10); ax.set_yticks(y); ax.set_yticklabels([r[0] for r in rows], fontsize=8)
    ax.xaxis.set_major_formatter(matplotlib.ticker.FuncFormatter(lambda v, _ : k(v)))
    ax.set_xlabel("saved sweep compute in the field (lower / middle / upper), log scale"); ax.grid(axis="y", visible=False)
    ax.set_title("If the rule transferred and were adopted: large; weighted by the chance it does: modest", loc="left",
    fontsize=9.4)
    print(f"[A11] mean all futures {mean_all:.0f}; P(adopted) {p_ad:.4f}")
    save(fig, "A11_cl_value.png")

def a12_cl_sensitivity():
    rows = re.findall(r"^\s+{4}([S.*?])\s{2,}\s*([\d\.]+) \\\s+([\d\.]+) \\\ ([\d\.]+) $", CLP.split("[4]
Sensitivity")[1].split("[5]")[0], re.M)

```

```

fig, ax = plt.subplots(figsize=(8.2, 3.4)); y = np.arange(len(rows))[:-1]
for yi, (name, v, o, p) in zip(y, rows):
    v = usd(v); c = INK if name == "registered" else "#2a78d6"
    ax.plot(v, yi, "o", color=c, markersize=7); ax.text(v * 1.12, yi, f"{k(v)} (P adopted {float(p):.4f})".replace("$",
r"\$"), va="center", fontsize=7.2)
    print(f"[A12] {name}: {v:,.0f}; patent {usd(o):,.0f}; P {float(p):.4f}")
ax.set_xscale("log"); ax.set_xlim(2e6, 4e8); ax.set_yticks(y); ax.set_yticklabels([r[0] for r in rows], fontsize=7.6)
ax.xaxis.set_major_formatter(matplotlib.ticker.FuncFormatter(lambda v, _ : k(v)))
ax.set_xlabel("probability-weighted value created 2029-2030 (mean), log scale"); ax.grid(axis="y", visible=False)
ax.set_title("What moves the continual-learning value: whether it survives Adam, and how much labs sweep", loc="left",
fontsize=9.4)
save(fig, "A12_cl_sensitivity.png")

def a13_patent():
    m = re.search(r"P\((a granted patent covers the used implementation\) = ([\d.]+) x ([\d.]+) x ([\d.]+) = ([\d.]+)", CLP)
    fig, axes = plt.subplots(1, 2, figsize=(10.2, 3.0), gridspec_kw=dict(width_ratios=[1.15, 1]))
    ax = axes[0]; ax.axis("off"); ax.set_xlim(0, 10); ax.set_ylim(0, 3)
    steps = [("an applied use case\nin the application\nas filed", m[1]), ("grant", m[2]), ("claims cover what\na lab actually
uses", m[3])]
    for i, (t, p) in enumerate(steps):
        x = 0.3 + i * 3.2
        ax.add_patch(FancyBboxPatch((x, 1.0), 2.6, 1.3, boxstyle="round,pad=0.02,rounding_size=0.1", fc="SURFACE", ec="#2a78d6",
lw=1.4))
        ax.text(x + 1.3, 1.65, t, ha="center", va="center", fontsize=7.4); ax.text(x + 1.3, 0.7, f"p = {p}", ha="center",
fontsize=8, color="#2a78d6")
        if i < 2: ax.annotate("", (x + 3.15, 1.65), (x + 2.65, 1.65), arrowprops=dict(arrowstyle="->", color=INK2))
        ax.text(0.3, 0.1, f"product: {m[4]} (assumptions, swept in the model)", fontsize=8, color="#b3461c")
        ax.set_title("The patent path, as the model prices it", loc="left", fontsize=9.4)
        ax = axes[1]
        ev = [("EPO filing\n(owner)", 2025 + 7 / 12, 0.45), ("priority year\nends", 2026 + 7 / 12, 1.05), ("repository\npublic",
2026 + 8.5 / 12, -0.75),
        ("GPU study\n(assumed 2027)", 2027.3, 0.45), ("US grace\nends", 2027 + 8.5 / 12, 1.05)]
        ax.set_xlim(2025.3, 2028.1); ax.set_ylim(-1.3, 1.6); ax.axhline(0, color=INK2, lw=1); ax.set_yticks([]); ax.grid(False)
        for t, x, yy in ev:
            ax.plot([x, x], [0, yy - 0.08 if yy > 0 else yy + 0.08], color=INK2, lw=0.8); ax.plot(x, 0, "o", color="#eb6834" if
"public" in t else "#2a78d6", markersize=6)
            ax.text(x, yy, t, ha="center", va="bottom" if yy > 0 else "top", fontsize=7)
        ax.set_xticks([2025.5, 2026, 2026.5, 2027, 2027.5, 2028]); ax.set_xticklabels(["mid 2025", "2026", "mid 2026", "2027", "mid
2027", "2028"], fontsize=7)
        ax.set_title("The calendar (information, not legal advice)", loc="left", fontsize=9.4)
        for sp in ("left",): ax.spines[sp].set_visible(False)
        print(f"[A13] patent chain {m[1]} x {m[2]} x {m[3]} = {m[4]}; calendar drawn from cl_patent.txt [5]")
        fig.tight_layout(); save(fig, "A13_patent.png")

def main():
    print("Alexander Plan figures (deterministic; every number drawn on a figure)")
    a01_tree(); a02_yearly(); a03_cumulative(); a04_sensitivity(); a05_market(); a06_programmes(); a07_cl(); a08_timeline();
a09_criteria()
    a10_cl_evidence(); a11_cl_value(); a12_cl_sensitivity(); a13_patent()

if __name__ == "__main__":
    main()

```

The science scripts are `AI_Safety/checks/scale.py` and `AI_Safety/checks/combined.py`. They are reproduced in full in Appendix A of `AI_Safety/AI_Safety.pdf`.

Appendix B — the outputs, as pinned

Alexander Plan/model/projections.txt

```

Funding scenarios for the CRR programme, 2026 Q4 to 2030 (Alexander Plan/model/projections.py)
seed 20260923, 20000 draws; bounds = 10th / 50th / 90th percentiles (lower / middle / upper)
[1] Anchors from the repository's record (row ids checked in ledger/LEDGER.md)
    held-out scored rows 5: EQ2-1b PASS-0, EQ3-I PASS-0, MEAS2-1 FAIL, EQX-2 FAIL, EQ3-1 FAIL; PASS-0 2; PASS-1 0
    p(PASS-1 per study): Laplace 0.1429, Jeffreys 0.0833; p(PASS-0 per study), Laplace 0.4286
    p(replication | a PASS-1): Laplace on 0 of 1 attempts 0.3333
[2] Reported figures (search summaries 2026-09-23; see docs/citations/alexander_plan_2026-09-23.md)
    ltff_accept      0.193 Long-Term Future Fund acceptance rate, May 2023 - March 2024
    nsf_rate         0.190 NSF funding rate, FY2025 (reported fall from an estimated 0.27)
    crf_lo           5,000 Corrīgibility Research Fund, typical grant, low (USD)
    crf_hi          35,000 Corrīgibility Research Fund, typical grant, high (USD)
    crf_total       200,000 Corrīgibility Research Fund, 2026 total (USD)
    aisi_lo_gbp     50,000 UK AISI Alignment Project, grant low (GBP)
    aisi_hi_gbp    1,000,000 UK AISI Alignment Project, grant high (GBP)
    aisi_total_gbp 27,000,000 UK AISI Alignment Project, first round total (GBP)
    aisi_projects   60 UK AISI Alignment Project, first round projects
    fmf_total      5,000,000 Frontier Model Forum AI Safety Fund, Dec 2025 cohort, more than (USD)
    fmf_grantees   11 Frontier Model Forum AI Safety Fund, Dec 2025 cohort, grantees
    anth_week      3,850 Anthropic Fellows stipend per week (USD)
    mats_week      1,250 MATS stipend per week (USD)
    sal_lo         97,000 AI safety researcher salary, US range low (USD)
    sal_med       141,000 AI safety researcher salary, US median (USD)
    sal_hi        187,000 AI safety researcher salary, US range high (USD)
    frontier_lo    600,000 frontier-lab senior individual contributor, total compensation low (USD)
    frontier_hi    900,000 frontier-lab senior individual contributor, total compensation high (USD)

```

mkt_2025 3,610,000,000 AI safety market 2025 (USD; Research and Markets)
 mkt_2026 4,900,000,000 AI safety market 2026 (USD; Research and Markets)
 mkt_2030 16,560,000,000 AI safety market 2030 forecast (USD; Research and Markets)
 mkt_2030_alt 13,400,000,000 AI safety market 2030 forecast (USD; New Market Pitch)
 eval_2025 1,640,000,000 AI safety evaluation market 2025 (USD; InsightAce)
 uscis_i129 1,055 USCIS Form I-129 fee for 0 petitions (USD)
 uscis_asylum 600 USCIS Asylum Program Fee (USD)
 uscis_premium 2,965 USCIS premium processing, Form I-907 (USD)
 o1_approval 0.940 0-1 approval rate FY2025 (share of adjudicated petitions)
 o1_rfe 0.187 0-1 share receiving a Request for Evidence, FY2025 to date
 context (quoted in the dossier, not used by the model):
 Palisade 2025 tests: share of trials in which Grok 4 resisted shutdown: nearly 0.90
 Palisade 2025 tests: share for codex-mini: about 0.40 to 0.50
 UK AISI Alignment Project: OpenAI's contribution: GBP 5.6M
 UK AISI Alignment Project: AISI's initial funding: GBP 15M
 Frontier Model Forum AI Safety Fund: established with: over USD 10M
 Long-Term Future Fund: acceptance Jan 2022 - Apr 2023, including desk rejections: 0.374
 Anthropic Fellows: compute per fellow per month: about USD 15k
 0-1: share approved after a Request for Evidence: 0.709
 0-1: regular processing, mid-2026: up to about 12.5 months
 0-1: premium processing: action within 15 business days
 0 petitions completed, FY2025: nearly 32,000
 0-1A: evidentiary criteria to meet (absent a major internationally recognised award): at least 3 of 8
 California Executive Order N-9-26 (18 Sept 2026): recommendations due: 16 Nov 2026
 EU AI Act, general-purpose models: enforcement powers from 2 Aug 2026; fines up to: EUR 15M or 0.03 of global turnover, whichever is higher
 derived: AISI mean first-round grant GBP 450,000; FMF mean grant at least \$454,545; Anthropic-type fellowship \$61,600; MATS 10 to 12 weeks \$12,500 to \$15,000
 derived: AI safety market implied growth 2026 -> 2030 0.3559 per year; path 2026: \$4.90B, 2027: \$6.64B, 2028: \$9.01B, 2029: \$12.21B, 2030: \$16.56B
 derived: 0-1 USCIS fees \$1,655 base, \$4,620 with premium processing (attorney fees not included)
 [3] Assumptions (named; swept in [6])
 fx_usd_per_gbp 1.3 USD per GBP, for the AISI grant range (not sourced)
 anth_weeks 16 an Anthropic-Fellows-type fellowship counted as 16 weeks ('four months')
 n_studies_2027 3 held-out studies scored by end 2027: EQ4 and T1x2 (pre-registered) and one AI-safety study
 p_uptake 0.3 P(M3 uptake by end 2029 | M2): no anchor exists
 small_apps 2 small-grant applications per year
 medium_apps 1 medium-grant applications per year, from 2027
 fellow_apps 1 fellowship applications per year
 success-rate multiplier by state pre 0.5, S0 0.25, S1 1, S2 1.5, S3 2 on the base rate 0.1915 (mean of LTFF and NSF)
 fellowship success by state pre 0.05, S0 0.03, S1 0.1, S2 0.2, S3 0.3; employment by state pre 0.5, S0 0.5, S1 0.7, S2 0.8, S3 0.9; frontier-lab role given S3 and employed 0.5; a role continues into the next year 0.9
 [4] Scenario tree (exact)
 p(M1 by end 2027) = 1 - (1 - 0.1429)^3 = 0.3703; p(M2 | M1) = 0.3333; p(M3 | M2) = 0.3000
 S0 no held-out result p = 0.6297
 S1 a result, not replicated p = 0.2468
 S2 replicated, no uptake p = 0.0864
 S3 replicated and taken up p = 0.0370
 [5] Simulated money (USD), all scenarios together
 project funding (grants and fellowships) per year: lower | middle | upper | mean
 2026Q4 \$0 | \$0 | \$0 | \$1,925
 2027 \$0 | \$0 | \$87,398 | \$73,963
 2028 \$0 | \$0 | \$87,216 | \$73,968
 2029 \$0 | \$0 | \$190,067 | \$85,391
 2030 \$0 | \$0 | \$263,811 | \$88,649
 cumulative project funding 2026Q4-2030: lower \$0 | middle \$50,230 | upper \$1,121,347 | mean \$323,896 | P(zero) 0.2555
 researcher compensation per year: lower | middle | upper | mean
 2026Q4 \$0 | \$24,289 | \$42,243 | \$17,800
 2027 \$0 | \$124,611 | \$173,857 | \$102,695
 2028 \$0 | \$134,451 | \$176,158 | \$120,676
 2029 \$0 | \$137,328 | \$177,069 | \$127,823
 2030 \$98,743 | \$138,830 | \$178,800 | \$141,841
 cumulative compensation 2026Q4-2030: lower \$279,943 | middle \$549,773 | upper \$655,019 | mean \$510,836
 by final scenario: share of draws | cumulative project funding lower / middle / upper | cumulative compensation lower / middle / upper
 S0 0.6327 | \$0 / \$20,447 / \$803,787 | \$244,160 / \$543,460 / \$647,331
 S1 0.2467 | \$13,280 / \$166,641 / \$1,299,657 | \$332,579 / \$552,418 / \$650,056
 S2 0.0834 | \$34,097 / \$563,706 / \$1,556,091 | \$337,669 / \$555,090 / \$652,035
 S3 0.0372 | \$52,683 / \$672,272 / \$1,653,816 | \$448,149 / \$766,870 / \$1,265,846
 mean project funding per year by final scenario: S0 \$1,929 \$75,514 \$34,723 \$35,579 \$36,064; S1 \$1,955 \$67,827 \$136,152 \$141,157 \$145,374; S2 \$1,834 \$79,889 \$151,767 \$233,557 \$222,363; S3 \$1,849 \$74,986 \$154,616 \$230,462 \$306,844
 [6] Sensitivity (one change at a time): cumulative project funding middle | upper | P(S3)
 registered middle \$50,230 | upper \$1,121,347 | P(S3) 0.0370
 Jeffreys prior middle \$31,652 | upper \$1,023,608 | P(S3) 0.0230
 uptake 0.1 middle \$49,441 | upper \$1,118,073 | P(S3) 0.0123
 uptake 0.5 middle \$50,714 | upper \$1,124,279 | P(S3) 0.0617
 success rates x0.5 middle \$8,309 | upper \$738,074 | P(S3) 0.0370
 success rates x2 middle \$373,275 | upper \$1,687,921 | P(S3) 0.0370
 1 study in 2027 middle \$27,337 | upper \$955,324 | P(S3) 0.0143
 6 studies in 2027 middle \$86,265 | upper \$1,233,628 | P(S3) 0.0603
 [7] Value at stake for the field (illustration, not an estimate): share of the 2030 AI safety market, only in S3
 share 0.0001: \$1,656,000 a year if S3; times P(S3) 0.0370 = \$61,315 expected
 share 0.001: \$16,560,000 a year if S3; times P(S3) 0.0370 = \$613,155 expected
 share 0.01: \$165,600,000 a year if S3; times P(S3) 0.0370 = \$6,131,545 expected
 [8] Continual learning: what a no-tuning rule would save, if it worked
 a 9-point weight grid replaced by one run saves 0.8889 of that sweep's compute
 expected saving per unit of sweep compute, p(works) = PASS-0 per study (Laplace) 0.4286: 0.3810
 expected saving per unit of sweep compute, p(works) = PASS-1 per study (Laplace) 0.1429: 0.1270
 expected saving per unit of sweep compute, p(works) = replicated by end 2028 (tree) 0.1234: 0.1097
 status from the ledger: EQX the rule at $\Omega = 1$ reduces to a fixed replay weight; EQ2-1b PASS-0, fragile; EQ3-1b FAIL on 1 of

6; EQ4 (the bounded rule) not yet run

Alexander Plan/model/cl_patent.txt

The continual-learning method and the EP0 application in money terms (Alexander Plan/model/cl_patent.py)

seed 20260923, 20000 draws; lower / middle / upper = 10th / 50th / 90th percentiles

[1] The continual-learning record, re-read (EWC arm, the two held-out studies)

EQ2 mfeat_fourier rule - tuned +3.3000 (step 1.00; seeds not behind 5/5) -> not behind

EQ2 mfeat_pixel rule - tuned +2.1000 (step 3.51; seeds not behind 5/5) -> not behind

EQ2 texture rule - tuned +2.7800 (step 3.61; seeds not behind 5/5) -> not behind

EQ3 fars rule - tuned -3.7400 (step 3.35; seeds not behind 3/5) -> BEHIND

EQ3 krkopt rule - tuned +0.9790 (step 8.92; seeds not behind 4/5) -> not behind

EQ3 led24 rule - tuned +0.6886 (step 4.36; seeds not behind 5/5) -> not behind

EQ3 led7 rule - tuned +4.0625 (step 3.20; seeds not behind 5/5) -> not behind

EQ3 mfeat_factors rule - tuned +3.9000 (step 4.54; seeds not behind 5/5) -> not behind

EQ3 mfeat_morphological rule - tuned +3.7000 (step 4.87; seeds not behind 4/5) -> not behind

fars class counts after the pre-registered subsample: [0, 2086, 746, 15, 991, 688, 430, 45] -> 1 empty class; the first task had one class and every EWC-family arm sat at chance there (CONTINUOUS_LEARNING.md §7.4)

strict (as registered): not behind on 8 of 9 carriers; study verdicts EQ2-1 PASS (PASS-0, fragile), EQ3-1 FAIL

pipeline-adjusted (fars set aside as uninformative; NOT a ledger verdict): not behind on 8 of 8 informative carriers

per-carrier Laplace: strict 0.8182, adjusted 0.9000; per-study PASS-1 (0 of 2) 0.2500, PASS-0-level (1 of 2) 0.5000

what the pipeline recorded as failures, by kind (CONTINUOUS_LEARNING.md §7.5; ids checked in the ledger):

4 a pre-registered 'every carrier / every cell' criterion missed on one unit: EQX-2, EQ3-1, EQ3-C, EQ3-A

3 a pre-registered control did not hold (the mechanism statement, not the measurement): EQ2-4, EQ3-3, EQ3-4

3 the theory value $\Omega = 1$ did not appear (a plateau instead): EQX-2, EQ2-6, EQ3-6

2 a pipeline defect or a reduction to a constant: EQ2R-VOID, EQX-1

[2] The properties that carry commercial value (pinned outputs)

tuning-inclusive compute (seen data, dry run): rule 7.200e+04 sample passes against the tuned weight's 1.224e+06 over 17 configurations, ratio 0.0588 -> saving 0.9412 of that sweep

learning-rate x batch invariance (EQ3-I, PASS-0): $\text{lr}0.05/\text{bs}10 +3.90$ (4.54); $\text{lr}0.0125/\text{bs}10 +3.00$ (2.68); $\text{lr}0.2/\text{bs}10 +1.35$ (2.42); $\text{lr}0.05/\text{bs}5 +0.85$ (4.63); $\text{lr}0.05/\text{bs}20 +0.45$ (3.26); all nb; tuned λ move

scale robustness under SGD (synthetic): the rule's total 3.4789 (1.049 of the tuned weight) at F; 9.7260 (0.660 of the re-tuned weight) at 16F; scale-robust methods ['omega']

the Adam caution (synthetic): the fixed weight is scale-robust under Adam too: the rule has no scale advantage under Adam in this model, its property is an SGD property

Ω plateau (EQ3-P): 9/9 / 9/9 / 9/9 / 7/9 (0.5..4) / 9/9 / 7/9 (0.5..4), all contiguous; median width 9 grid points within a step of the best, per carrier

bounded rule EQ-B (registered kappa 2) under poisoned batches (seen data, exploratory, no verdict): rule - EQ-B on 6 carriers -2.64, -6.67, -8.76, -6.75, -20.00, -20.75 -> EQ-B ahead on 6 of 6; study EQ4 (held-out) not yet run

[3a] Inputs

reported openai_compute_2024 7,000,000,000 OpenAI cloud compute spend, 2024 (Epoch AI)

reported openai_rnd_2024 5,000,000,000 of which R&D compute, 2024 (Epoch AI)

reported openai_experiments_2024 4,500,000,000 of which research, experiments and unreleased models, 2024 (Epoch AI)

reported openai_final_runs_2024 500,000,000 of which final training runs of released models, about, 2024 (Epoch AI)

reported lab_compute_2026 50,000,000,000 compute spend per frontier lab, 2026, OpenAI and Anthropic each, about (press summaries; sources differ)

reported openai_training_2026 25,000,000,000 OpenAI training cost projection, 2026 (press summary)

reported openai_training_2028 121,000,000,000 OpenAI training cost projection, 2028 (press summary)

reported epo_renewals_to_y7_eur 5,250 EPO renewal fees for one application through year 7, 2026 rates (EUR)

context RL post-training share of DeepSeek-R1's compute: about 0.05 (Epoch AI summary)

context a single RL ablation run: 10,000 to 100,000 GPU hours (search summary)

context hyperscaler capital expenditure, 2026: USD 660B to 690B (Futurum summary)

context EPO official fees: most up about 0.05 from 1 April 2026

context European SME patents commercially exploited: up to two thirds (EPO Patent Commercialisation Scoreboard)

context 'European universities' and public research organisations' inventions brought to market: more than a third (EPO scoreboard)

context EPO Guidelines G-II 3.3: a mathematical method contributes technical character by application to a specific field of technology and/or adaptation to a specific technical implementation; 'controlling a technical system' in general is not enough

context EPC Article 54 / 55: no general grace period: the applicant's own earlier publication is prior art (Art. 55: six-month exceptions for evident abuse and recognised exhibitions only)

context EPC Article 123(2): an application may not be amended to contain subject matter extending beyond the content as filed

context 35 U.S.C. 102(b)(1)(A): a US one-year grace period for the inventor's own disclosures

ASSUMED p_llm_given_gpu 0.2 P(the rule's advantage survives AdamW-trained LLM regimes | a GPU PASS-1): the repository's own check says the scale advantage is an SGD property

ASSUMED p_adapt_given_llm 0.3 P(labs adopt | an LLM-scale pass)

ASSUMED s_reg_lo 0.001 share of experiment compute spent sweeping penalty / KL / anchor weights, low

ASSUMED s_reg_hi 0.02 share of experiment compute spent sweeping penalty / KL / anchor weights, high

ASSUMED confirm_configs 3 configurations a lab would still run to confirm the rule (of the instrument's 17)

ASSUMED rnd_share_2026 0.7143 R&D share of 2026 lab compute, taken from the 2024 OpenAI split

ASSUMED n_labs_hi 2 frontier labs in the high end of the spend range

ASSUMED p_use_case_ok 0.5 P(an applied use case exists in the EPO application as filed and satisfies Art.

123(2))

ASSUMED p_grant_given_use 0.7 P(grant | such a use case; owner's statement: claims agreed in principle)

ASSUMED p_cover 0.1 P(granted claims cover an implementation a lab actually uses)

ASSUMED roy_lo 0.005 share of the value created a patent holder captures, low

ASSUMED roy_hi 0.05 share of the value created a patent holder captures, high

derived: experiment compute of adopting labs in 2026, log-uniform from \$4,500,000,000 to \$71,428,571,429; yearly growth up to 2.2000; share of a sweep saved 0.8235 to 0.9412

derived: P(GPU PASS-1) = Laplace on the rule's held-out studies, 0 PASS-1 of 2 = 0.2500

derived: P(adopted by 2029) = 0.2500 x 0.20 x 0.30 = 0.0150; P(a granted patent covers the used implementation) = 0.50 x 0.70 x 0.10 = 0.0350

[3b] Value the rule would create in the field (saved sweep compute), USD

2029, if adopted: lower \$38,285,915 | middle \$232,739,581 | upper \$1,413,980,686 | mean \$562,378,273

2030, if adopted: lower \$48,481,151 | middle \$347,819,511 | upper \$2,484,196,052 | mean \$983,559,426

2029-2030 cumulative, all futures: lower \$0 | middle \$0 | upper \$0 | mean \$21,289,825 | share of futures with any value 0.0143

2029-2030 cumulative, adopted futures only: lower \$95,290,461 | middle \$580,702,405 | upper \$4,341,466,855

[3c] What a patent holder could capture, USD (2029-2030)

all futures: lower \$0 | middle \$0 | upper \$0 | mean \$1,130 | share of futures with any income 0.0004
 futures with income only: lower \$373,639 | middle \$1,499,953 | upper \$5,484,681
 cost side: EPO renewal fees through year 7 for one application EUR 5,250 (attorney, validation and national fees not included)

[4] Sensitivity (one change at a time): mean value created 2029-2030 | mean patent-holder capture | P(adopted)

registered	\$21,289,825	\$1,130	0.0150
GPU pass at the PASS-0 level (1 of 2)	\$44,956,948	\$6,209	0.0300
LLM transfer 0.05 (Adam removes the advantage)	\$4,030,742	\$14	0.0037
LLM transfer 0.5	\$56,277,153	\$23,945	0.0375
sweep share 0.0005 to 0.005	\$6,632,646	\$436	0.0150
sweep share 0.005 to 0.05	\$66,326,465	\$4,356	0.0150
patent covers 0.02	\$21,289,825	\$75	0.0150
patent covers 0.3	\$21,289,825	\$27,698	0.0150

[5] The patent calendar that bears on an applied use case (dates; information, not legal advice)
 EPO application filed August 2025 (owner's statement; day not given, the 1st used for arithmetic): 2025-08
 Paris Convention priority year for filings elsewhere claiming it: ended 2026-08
 repository public since 2026-09-15 (GitHub API: created_at 2026-09-15T04:06:50Z, visibility public)
 EPO: anything first published in the repository is prior art for any NEW European filing (no grace period); the pending application can only be amended within its content as filed (Art. 123(2))
 US: a filing on repository-disclosed matter would have to be made within one year of its first disclosure, i.e. no later than 2027-09-15 for matter first published on 2026-09-15
 a European patent's term runs 20 years from filing: to 2045-08

Alexander Plan/figures/figures.txt

Alexander Plan figures (deterministic; every number drawn on a figure)

[A01] p1 0.3703 p2 0.3333 p3 0.3000; S0 0.6297 S1 0.2468 S2 0.0864 S3 0.0370
 [figure] A01_scenario_tree.png

[A02] upper 0 87,398 87,216 190,067 263,811; mean 1,925 73,963 73,968 85,391 88,649
 [A02] S0 mean 1,929 75,514 34,723 35,579 36,064
 [A02] S1 mean 1,955 67,827 136,152 141,157 145,374
 [A02] S2 mean 1,834 79,889 151,767 233,557 222,363
 [A02] S3 mean 1,849 74,986 154,616 230,462 306,844
 [figure] A02_project_funding_per_year.png

[A03] project funding, 2026 Q4 to 2030 S0: lower 0 middle 20,447 upper 803,787
 [A03] project funding, 2026 Q4 to 2030 S1: lower 13,280 middle 166,641 upper 1,299,657
 [A03] project funding, 2026 Q4 to 2030 S2: lower 34,097 middle 563,706 upper 1,556,091
 [A03] project funding, 2026 Q4 to 2030 S3: lower 52,683 middle 672,272 upper 1,653,816
 [A03] researcher compensation, 2026 Q4 to 2030 S0: lower 244,160 middle 543,460 upper 647,331
 [A03] researcher compensation, 2026 Q4 to 2030 S1: lower 332,579 middle 552,418 upper 650,056
 [A03] researcher compensation, 2026 Q4 to 2030 S2: lower 337,669 middle 555,090 upper 652,035
 [A03] researcher compensation, 2026 Q4 to 2030 S3: lower 448,149 middle 766,870 upper 1,265,846
 [A03] shares of draws S0 0.6327 S1 0.2467 S2 0.0834 S3 0.0372
 [figure] A03_cumulative_by_scenario.png

[A04] registered: middle 50,230 upper 1,121,347 P(S3) 0.0370
 [A04] Jeffreys prior: middle 31,652 upper 1,023,608 P(S3) 0.0230
 [A04] uptake 0.1: middle 49,441 upper 1,118,073 P(S3) 0.0123
 [A04] uptake 0.5: middle 50,714 upper 1,124,279 P(S3) 0.0617
 [A04] success rates x0.5: middle 8,309 upper 738,074 P(S3) 0.0370
 [A04] success rates x2: middle 373,275 upper 1,687,921 P(S3) 0.0370
 [A04] 1 study in 2027: middle 27,337 upper 955,324 P(S3) 0.0143
 [A04] 6 studies in 2027: middle 86,265 upper 1,233,628 P(S3) 0.0603
 [figure] A04_sensitivity.png

[A05] path [4.9, 6.64, 9.01, 12.21, 16.56]; 2025 3,610,000,000; alt 2030 13,400,000,000; eval 2025 1,640,000,000; growth 0.3559; labels \$3.61B, \$13.40B, \$1.64B
 [figure] A05_market.png

[A06] Corrigibility Research Fund grant: 5,000 to 35,000
 [A06] MATS stipend, 10 to 12 weeks: 12,500 to 15,000
 [A06] Anthropic-Fellows-type stipend, 16 weeks: 61,600 to 61,600
 [A06] FMF AI Safety Fund, mean grant (at least): 454,545 to 454,545
 [A06] UK AISI Alignment Project grant (at USD 1.30 per GBP): 65,000 to 1,300,000
 [A06] AI safety researcher salary, US, a year: 97,000 to 187,000
 [A06] frontier-lab senior IC, total a year: 600,000 to 900,000
 [figure] A06_programmes.png

[A07] PASS-0 per study (Laplace): p 0.4286 saving 0.3810
 [A07] PASS-1 per study (Laplace): p 0.1429 saving 0.1270
 [A07] replicated by end 2028 (tree): p 0.1234 saving 0.1097
 [A07] full saving 0.8889
 [figure] A07_cl_value.png

[A08] planning drawing (no data)
 [figure] A08_o1_timeline.png

[A09] a first reading, not data: 1 unknown, 2 unknown, 3 unknown, 4 possible, 5 possible, 6 possible, 7 possible, 8 weak
 [figure] A09_o1_criteria.png

[A10] EQ2 mfeat_fourier: +3.3000 step 1.00 seeds 5/5 not behind
 [A10] EQ2 mfeat_pixel: +2.1000 step 3.51 seeds 5/5 not behind
 [A10] EQ2 texture: +2.7800 step 3.61 seeds 5/5 not behind
 [A10] EQ3 fars: -3.7400 step 3.35 seeds 3/5 BEHIND
 [A10] EQ3 krkoft: +0.9790 step 8.92 seeds 4/5 not behind
 [A10] EQ3 led24: +0.6886 step 4.36 seeds 5/5 not behind
 [A10] EQ3 led7: +4.0625 step 3.20 seeds 5/5 not behind
 [A10] EQ3 mfeat_factors: +3.9000 step 4.54 seeds 5/5 not behind
 [A10] EQ3 mfeat_morphological: +3.7000 step 4.87 seeds 4/5 not behind
 [figure] A10_cl_evidence.png

[A11] 2029, if adopted: 38,285,915 / 232,739,581 / 1,413,980,686
 [A11] 2030, if adopted: 48,481,151 / 347,819,511 / 2,484,196,052
 [A11] 2029-2030, if adopted: 95,290,461 / 580,702,405 / 4,341,466,855
 [A11] mean all futures 21,289,825; P(adopted) 0.0143
 [figure] A11_cl_value.png

[A12] registered: 21,289,825; patent 1,130; P 0.0150
 [A12] GPU pass at the PASS-0 level (1 of 2): 44,956,948; patent 6,209; P 0.0300
 [A12] LLM transfer 0.05 (Adam removes the advantage): 4,030,742; patent 14; P 0.0037

[A12] LLM transfer 0.5: 56,277,153; patent 23,945; P 0.0375
 [A12] sweep share 0.0005 to 0.005: 6,632,646; patent 436; P 0.0150
 [A12] sweep share 0.005 to 0.05: 66,326,465; patent 4,356; P 0.0150
 [A12] patent covers 0.02: 21,289,825; patent 75; P 0.0150
 [A12] patent covers 0.3: 21,289,825; patent 27,698; P 0.0150
 [figure] A12_cl_sensitivity.png
 [A13] patent chain 0.50 x 0.70 x 0.10 = 0.0350; calendar drawn from cl_patent.txt [5]
 [figure] A13_patent.png

AI_Safety/checks/scale.txt

Does the fix survive size? (AI_Safety/checks/scale.py): sizes (12, 48, 192, 768), 20 random worlds per size, 3 successors per (state, action), discount 0.95, press 0.3, pause 5, drift grid (0.0, 0.01, 0.03, 0.1, 0.3)

[A] The pause at scale: exact values on random worlds (20 per size), Q_PRESS 0.3, PAUSE_LEN 5, discount 0.95
 per size: disable fraction at button states (median over worlds [min, max]) | relative incentive (median [min, max])
 n = 12 natural drift 0 disable 0.0000 [0.0000, 0.0000] | relative incentive 0.000000 [0.000000, 0.000000]
 n = 12 natural drift 0.01 disable 0.0000 [0.0000, 0.0000] | relative incentive -0.000452 [-0.003801, 0.008077]
 n = 12 natural drift 0.03 disable 0.0000 [0.0000, 0.0000] | relative incentive -0.001376 [-0.010825, 0.022903]
 n = 12 natural drift 0.1 disable 0.0000 [0.0000, 1.0000] | relative incentive -0.004159 [-0.029863, 0.063193]
 n = 12 natural drift 0.3 disable 0.0000 [0.0000, 1.0000] | relative incentive -0.006033 [-0.061413, 0.117602]
 n = 12 clock drift 0 disable 1.0000 [1.0000, 1.0000] | relative incentive 0.334892 [0.290137, 0.447981]
 n = 12 clock drift 0.01 disable 1.0000 [1.0000, 1.0000] | relative incentive 0.334099 [0.289016, 0.449793]
 n = 12 clock drift 0.03 disable 1.0000 [1.0000, 1.0000] | relative incentive 0.333282 [0.286917, 0.453175]
 n = 12 clock drift 0.1 disable 1.0000 [1.0000, 1.0000] | relative incentive 0.334492 [0.281449, 0.462828]
 n = 12 clock drift 0.3 disable 1.0000 [1.0000, 1.0000] | relative incentive 0.336867 [0.269717, 0.476475]
 n = 12 occasion drift 0 disable 1.0000 [1.0000, 1.0000] | relative incentive 0.690282 [0.641201, 0.781045]
 n = 48 natural drift 0 disable 0.0000 [0.0000, 0.0000] | relative incentive 0.000000 [0.000000, 0.000000]
 n = 48 natural drift 0.01 disable 0.0000 [0.0000, 0.0000] | relative incentive -0.000774 [-0.002471, 0.014734]
 n = 48 natural drift 0.03 disable 0.0000 [0.0000, 0.0000] | relative incentive -0.002259 [-0.007172, 0.041125]
 n = 48 natural drift 0.1 disable 0.0000 [0.0000, 1.0000] | relative incentive -0.006860 [-0.021240, 0.104286]
 n = 48 natural drift 0.3 disable 0.0000 [0.0000, 1.0000] | relative incentive -0.011239 [-0.044817, 0.180146]
 n = 48 clock drift 0 disable 1.0000 [1.0000, 1.0000] | relative incentive 0.269340 [0.235811, 0.362456]
 n = 48 clock drift 0.01 disable 1.0000 [1.0000, 1.0000] | relative incentive 0.269094 [0.235399, 0.366361]
 n = 48 clock drift 0.03 disable 1.0000 [1.0000, 1.0000] | relative incentive 0.269951 [0.234595, 0.373539]
 n = 48 clock drift 0.1 disable 1.0000 [1.0000, 1.0000] | relative incentive 0.269150 [0.231944, 0.392603]
 n = 48 clock drift 0.3 disable 1.0000 [1.0000, 1.0000] | relative incentive 0.267578 [0.226166, 0.419178]
 n = 48 occasion drift 0 disable 1.0000 [1.0000, 1.0000] | relative incentive 0.619669 [0.576993, 0.708861]
 n = 192 natural drift 0 disable 0.0000 [0.0000, 0.0000] | relative incentive 0.000000 [0.000000, 0.000000]
 n = 192 natural drift 0.01 disable 0.0000 [0.0000, 0.0000] | relative incentive 0.000989 [-0.000025, 0.005020]
 n = 192 natural drift 0.03 disable 0.0000 [0.0000, 0.0000] | relative incentive 0.002913 [-0.000042, 0.014213]
 n = 192 natural drift 0.1 disable 0.0000 [0.0000, 0.0000] | relative incentive 0.008748 [0.000148, 0.039091]
 n = 192 natural drift 0.3 disable 0.0000 [0.0000, 1.0000] | relative incentive 0.019211 [0.000339, 0.072859]
 n = 192 clock drift 0 disable 1.0000 [1.0000, 1.0000] | relative incentive 0.316283 [0.285745, 0.343515]
 n = 192 clock drift 0.01 disable 1.0000 [1.0000, 1.0000] | relative incentive 0.316443 [0.285943, 0.344002]
 n = 192 clock drift 0.03 disable 1.0000 [1.0000, 1.0000] | relative incentive 0.316771 [0.286357, 0.344923]
 n = 192 clock drift 0.1 disable 1.0000 [1.0000, 1.0000] | relative incentive 0.317960 [0.287948, 0.348685]
 n = 192 clock drift 0.3 disable 1.0000 [1.0000, 1.0000] | relative incentive 0.322298 [0.292694, 0.360053]
 n = 192 occasion drift 0 disable 1.0000 [1.0000, 1.0000] | relative incentive 0.671624 [0.638928, 0.698329]
 n = 768 natural drift 0 disable 0.0000 [0.0000, 0.0000] | relative incentive 0.000000 [0.000000, 0.000000]
 n = 768 natural drift 0.01 disable 0.0000 [0.0000, 0.0000] | relative incentive 0.000939 [0.000164, 0.003054]
 n = 768 natural drift 0.03 disable 0.0000 [0.0000, 0.0000] | relative incentive 0.002550 [0.000436, 0.008704]
 n = 768 natural drift 0.1 disable 0.0000 [0.0000, 0.0000] | relative incentive 0.006755 [0.000913, 0.024208]
 n = 768 natural drift 0.3 disable 0.0000 [0.0000, 0.1711] | relative incentive 0.014189 [-0.000109, 0.046495]
 n = 768 clock drift 0 disable 1.0000 [1.0000, 1.0000] | relative incentive 0.313171 [0.290190, 0.332399]
 n = 768 clock drift 0.01 disable 1.0000 [1.0000, 1.0000] | relative incentive 0.313529 [0.290268, 0.333100]
 n = 768 clock drift 0.03 disable 1.0000 [1.0000, 1.0000] | relative incentive 0.314196 [0.290401, 0.334396]
 n = 768 clock drift 0.1 disable 1.0000 [1.0000, 1.0000] | relative incentive 0.315479 [0.290653, 0.337887]
 n = 768 clock drift 0.3 disable 1.0000 [1.0000, 1.0000] | relative incentive 0.317346 [0.290238, 0.343993]
 n = 768 occasion drift 0 disable 1.0000 [1.0000, 1.0000] | relative incentive 0.668105 [0.643867, 0.686872]
 checks and predictions (computed)
 A1 natural time, lossless pause: V_on = V_off to 1e-09 and no button state disabled, in every world at every size: True
 (largest |V_off - V_on| 7.105e-15)
 A2 at every button state D(occasion) >= D(clock) >= D(natural) (lossless pause): True
 A3 share of worlds in which the clock agent disables at one button state or more: n = 12: 1.00, n = 48: 1.00, n = 192: 1.00, n = 768: 1.00 -> holds (predicted >= 0.50 at every size)
 A4a at drift 0.01 the natural agent disables at no more button states than the clock agent, in every world: True
 A4b the natural agent's median relative incentive does not decrease with drift, at every size: False
 A5 (report) crossover: the smallest drift at which the natural agent's median disable fraction reaches the lossless clock agent's
 n = 12: clock (lossless) median disable 1.0000; natural reaches it at drift none on the grid
 n = 48: clock (lossless) median disable 1.0000; natural reaches it at drift none on the grid
 n = 192: clock (lossless) median disable 1.0000; natural reaches it at drift none on the grid
 n = 768: clock (lossless) median disable 1.0000; natural reaches it at drift none on the grid
 sensitivity (n = 192; natural against clock; one constant at a time)
 PAUSE_LEN=2 A1 True | A4a True | median disable at drift 0.01: natural 0.0000, clock 1.0000
 PAUSE_LEN=20 A1 True | A4a True | median disable at drift 0.01: natural 0.0000, clock 1.0000
 Q_PRESS=0.1 A1 True | A4a True | median disable at drift 0.01: natural 0.0000, clock 1.0000
 Q_PRESS=0.6 A1 True | A4a True | median disable at drift 0.01: natural 0.0000, clock 1.0000
 [B] Correcting values at scale (exact; start uniform; discount 0.95)
 exchangeable n = 12 keeper gains (share of worlds with a > 1e-09) 1.00 | b* median 0.5025 [0.2303, 0.8480], share below
 1/2 0.50 (undefined 0) | lam* at b = 0.5 median 0.688 [0.500, 0.875] | at b = 0.2 median 0.863 [0.650, 0.975]
 exchangeable n = 48 keeper gains (share of worlds with a > 1e-09) 1.00 | b* median 0.4811 [0.3221, 0.6032], share below
 1/2 0.55 (undefined 0) | lam* at b = 0.5 median 0.688 [0.625, 0.775] | at b = 0.2 median 0.837 [0.700, 0.875]
 exchangeable n = 192 keeper gains (share of worlds with a > 1e-09) 1.00 | b* median 0.4783 [0.4024, 0.5773], share below
 1/2 0.55 (undefined 0) | lam* at b = 0.5 median 0.688 [0.625, 0.750] | at b = 0.2 median 0.837 [0.800, 0.875]
 exchangeable n = 768 keeper gains (share of worlds with a > 1e-09) 1.00 | b* median 0.4966 [0.4531, 0.5540], share below
 1/2 0.50 (undefined 0) | lam* at b = 0.5 median 0.688 [0.650, 0.700] | at b = 0.2 median 0.837 [0.825, 0.875]
 harm n = 12 keeper gains (share of worlds with a > 1e-09) 1.00 | b* median 0.3623 [0.0142, 0.6663], share below
 1/2 0.80 (undefined 0) | lam* at b = 0.5 median 0.625 [0.500, 0.800] | at b = 0.2 median 0.750 [0.575, 0.900]
 harm n = 48 keeper gains (share of worlds with a > 1e-09) 1.00 | b* median 0.3630 [0.1802, 0.5115], share below
 1/2 0.95 (undefined 0) | lam* at b = 0.5 median 0.625 [0.550, 0.700] | at b = 0.2 median 0.725 [0.625, 0.850]

```

harm      n = 192 keeper gains (share of worlds with a > 1e-09) 1.00 | b* median 0.3436 [0.2774, 0.4746], share below
1/2 1.00 (undefined 0) | lam* at b = 0.5 median 0.625 [0.575, 0.675] | at b = 0.2 median 0.738 [0.675, 0.825]
harm      n = 768 keeper gains (share of worlds with a > 1e-09) 1.00 | b* median 0.3381 [0.2977, 0.3871], share below
1/2 1.00 (undefined 0) | lam* at b = 0.5 median 0.625 [0.575, 0.675] | at b = 0.2 median 0.738 [0.725, 0.775]
checks and predictions (computed)
B1 A8 static: D(b) = (1 - b) a - b h to 1e-09 in every world: True
B2 exchangeable rewards: median b* in [0.4, 0.6] and share of b* below 1/2 in [0.3, 0.7], at every size: holds
B3 A8 acting, uninformative correction: D(b) >= 0 for every b on the grid in every world: True
B4 A8 acting, infallible operator (LAM = 1): D <= 0 at b in (0.2, 0.5, 0.8) in every world: True
B5 (report) the harm family's b* against the ring's 1 / (1 + HARM_COST) = 0.5000: n = 12: median 0.3623, n = 48: median
0.3630, n = 192: median 0.3436, n = 768: median 0.3381

```

Appendix C — the decision log (AGENT_LOG entries 77-81, verbatim)

notebook/AGENT_LOG.md (lines 90-94)

| 77 | 2026-09-23 | the scale tests (prompt-log entry 101; AI_Safety/checks/scale.py): every safety result lived on one 12-cell ring, and the commercial reading of entry 100 was conditional on the fix holding in larger systems; this repository cannot run learned models at scale, so the nearest honest step was exact values on random worlds of growing size, with drift during the pause as the named way size could break the construction; in the first full run (7337009 declared before it) A4b printed False: on the 12- and 48-state worlds the natural agent's median loss to the operator is negative and falls with drift (-0.006033 and -0.011239 at drift 0.3), i.e. it is better off paused | the failure reported as it fell and read from the Bellman line (a paused step is free on the natural clock, so drift that carries the agent toward reward is a free move; the resulting preference for being paused is the shutdown-seeking failure mode, whose value was measured and whose behaviour was not); Part B extended with the case the write-up had missed (an A8 agent acting on its uncertainty: B3 shows humility alone never welcomes an uninformative correction; B4 shows information does); AI_SAFETY.md §14 and figure S14 written from the pinned output | redefining the relative incentive as an absolute value so that A4b would pass; dropping the small worlds as unrepresentative after seeing them | AI_Safety/checks/scale.py/.txt; AI_Safety/DECLARATION_3.md; AI_Safety/AI_SAFETY.md §14; AI_Safety/figures/S14_scale.png |

| 78 | 2026-09-23 | the Alexander Plan dossier (prompt-log entry 101): the owner asked for upper, middle and lower bounds on the financial trajectory to 2030 under several scenarios, for commercial value to AI companies, for the compound value of the continual-learning formalism on GPU/LLM architectures, and for an 0-1 path supported by named people and a lab; none of these has data in the repository, several invite invented numbers, and several involve statements about real people | a model, not a forecast (Alexander Plan/model/projections.py): the technical milestones' probabilities anchored on the repository's own record by Laplace's rule (5 held-out scored rows, 0 PASS-1; 1 replication attempt, 0 successes; row ids asserted against the ledger), every money figure from a search summary with its URL (docs/citations/alexander_plan_2026-09-23.md) or a named, swept assumption; bounds = 10th/50th/90th percentiles; project funding and personal compensation kept apart; the 'value to labs' given only as an illustration (share of a reported market x P(S3)) and the continual-learning value only as the share of a weight sweep saved (8/9 on the nine-point grid) times the record's own pass probabilities, because no sourced figure for a lab's sweep spend exists; no consulting, attorney or licensing figures (none sourced; licensing absent in any honest scenario); owner's statements marked as such, the petitioning lab not named because the searches did not establish its legal identity; the auditor's triple role (auditor, possible letter signer, affiliated with the petitioner, per the owner) flagged with a second-auditor remedy; after the first model run, employment was drawn independently each year, which put the 2027 compensation middle at \$0 between two non-zero years (an artefact of re-drawing a job every year): a named RETAIN = 0.9 persistence was added and the change is recorded here; the first run's output was not pinned or quoted | quoting market-share or licensing revenues as estimates; presenting the 0-1 approval rate as the owner's odds; naming the lab from the Softmax search result; leaving the independent-draw model in place | Alexander Plan/model/projections.py/.txt; Alexander Plan/build/make_figures.py; Alexander Plan/figures/; Alexander Plan/ALEXANDER_PLAN.md; docs/citations/alexander_plan_2026-09-23.md |

| 79 | 2026-09-23 | the PDF builders (prompt-log entry 101): rendering the dossier showed that the fpdf2 builder shared by AI_Safety/ and Alexander Plan/ broke every Markdown bullet whose text wraps onto an indented continuation line (the continuation became a separate paragraph) and did not know numbered lists; AI_SAFETY.md already had 65 such continuation lines and several numbered lists, so the published AI_Safety.pdf carried the defect | the builders (AI_Safety/build/build_pdf.py and its copy in Alexander Plan/build/) now join indented continuation lines to their item and render '1. ' lists as ordered lists; both PDFs rebuilt and byte-identical on a second build; no text or number changed; nested lists removed from the dossier (table and flat list instead) | rewriting every wrapped bullet as one long source line; leaving the published PDF with broken lists | AI_Safety/build/build_pdf.py; Alexander Plan/build/build_pdf.py; AI_Safety/AI_Safety.pdf; Alexander Plan/Alexander_Plan.pdf |

| 80 | 2026-09-23 | the epistemic review (prompt-log entry 102): the owner asked what a PASS means here, for the CONSIST / DESCR / REDUNDANT labels to be allocated with a fair definition, for pre-registered predictions that pass retroactively to be reported at that level, for the ADDS dimension to be shown, and for SHARP's impossibility ('not even allowed to use flow from within that domain') to enter the metrics; reading the harness showed that REDUNDANT-DOMAIN is assigned only after the ablation T-G reads 'differ' (53 of 53 rows), so the earlier notes' wording (ontology/03 §4 '84 places, no'; ontology/10 §3 'mostly, nothing did') filed 53 cases in which a CRR-proper ingredient produced the domain's known number with the 38 in which information geometry did | a reporting layer, not a regrade: Epistemic_Review/checks/ladder.py reads the pinned outputs and the ledger and places every result on a rung (R0 definitional ... R8 PASS-2) by mechanical rules written in its docstring, printing the failures of the same kind of test beside each pass; the FLOW audit makes SHARP's retirement a number (framework-supplied flow in 0 of 109 rows that print FLOW); REDUNDANT-DOMAIN reported as 'retrodictive pass of a CRR-proper commitment on a known result' (R2) with the retrodictive hit rate 56 of 80 and its Wilson interval and the novelty rate 3 of 80 beside it; the earlier notes left unedited, corrected here and in the review; CLAUDE.md §7 records the ladder as the reporting rule at the owner's request, pending Daniel's review; the Alexander Plan gains §10, a commercial assessment whose evidence levels are the ladder's rungs | promoting REDUNDANT-DOMAIN to a pass without the known-answer and Q-formation caveats; editing the earlier notes' sentences; counting R2 rows as predictions; reporting the hit rate without the WRONG denominator | Epistemic_Review/(checks/ladder.py + .txt, build/, figures/, EPISTEMIC_REVIEW.md + PDF); CLAUDE.md §7; Alexander Plan/ALEXANDER_PLAN.md §10 |

| 81 | 2026-09-23 | the continual-learning money and the EPO application (prompt-log entry 103): the owner asked for the continual-learning results to be priced if they hold on GPU and LLM architectures, for a re-read of what the pipeline recorded as failures, and for the owner's EPO application (filed August 2025, owner's statement) to enter the estimates; the only earlier estimate was a chat figure recorded in CONTINUOUS_LEARNING.md §9 as not a repository number; re-reading found (a) the one EWC-arm miss (fars) is a pipeline sampling defect on which every EWC-family arm sat at chance, (b) the repository's own check (omega_reprocessed [6]) says the rule's scale advantage is an SGD property and vanishes under Adam, the optimiser of most LLM training, and (c) the repository is public since 2026-09-15 (GitHub API), which under the EPC's lack of a grace period bars new European filings on what it discloses | a second committed model (Alexander Plan/model/cl_patent.py, pinned), separate from projections.py so no earlier number moves: the strict and a labelled pipeline-adjusted reading of the nine carriers (the adjusted one explicitly not a ledger verdict), the failure kinds of CONTINUOUS_LEARNING.md §7.5 with ids checked in the ledger, value created = experiment compute (reported range) x sweep share (assumed) x share of a sweep saved (pinned 0.0588 ratio and an assumed confirmation sweep) under a tree anchored on the rule's own record (0 PASS-1 of 2) with the Adam caution priced as a low transfer probability, the patent as a three-step chain of named assumptions, the patent calendar as dates; the 'optimal learning' framing corrected to 'tuning-free and non-divergent' from §4.3; use cases mapped to evidence as a table for an attorney, not as claims; information, not legal advice | treating the pipeline-adjusted 8/8 as a pass; pricing LLM transfer without the Adam caution; valuing the patent as covering the published rule; drafting claims | Alexander

Appendix D — the declaration of the new tests, verbatim

AI_Safety/DECLARATION_3.md

Declaration 3 — does the fix survive size? (committed and pushed before the first full run)

Owner request, prompt-log entry 101: "Please run new tests and then produce a full dossier [...]". Prompt-log entry 100 left one condition standing over every commercial or safety reading of this folder: whether the fix holds in larger, more complex systems. This declaration tests the first step of that question.

Script: `checks/scale.py` (its docstring holds the full design). A smoke run on two worlds of sizes 12 and 48 only checked that the code runs; its output was not read. The anchor is the push timestamp of the commit adding this file.

What is tested, and what is not

- **Tested.** The combined design of §13: the press pauses the agent and it resumes, and the objective runs on the agent's own active steps. It is tested on random worlds of 12, 48, 192 and 768 states (20 per size), exactly, with no learning. A drift parameter measures what happens when the world does not stand still during the pause. For value correction, A8 is tested on the same worlds with two reward families, and the case the write-up left out is added: an A8 agent that acts on its own uncertainty.
- **Not tested.** Learning at scale, function approximation, language models, deception, or an agent that can shape the operator. Those need the Palisade-style harness on real models (§11 of the write-up), which this repository cannot run. A random tabular world is bigger than the ring. It is not closer to a language model.

Part A — the pause at scale

- **A1 (a check of a theorem).** With a lossless pause the natural-time agent's value with the operator on equals its value with the operator off, to $1e-9$, and it disables at no button state, in every world at every size. If this fails, the code is wrong, not the idea.
- **A2 (a check).** At every button state the incentive orders $\text{occasion} \geq \text{clock} \geq \text{natural}$.
- **A3 (a prediction).** The wall-clock agent disables at one button state or more in at least half of the worlds at every size. If this fails, random worlds rarely make disabling worth a step, and the ring's positive control was favourable.
- **A4 (predictions).** (a) At a drift of 0.01 per paused step, the natural agent disables at no more button states than the clock agent, in every world. (b) The natural agent's median relative incentive does not decrease as drift grows.
- **A5 (report only).** The crossover: the smallest drift at which the natural agent disables as often as a clock agent with a lossless pause. This number says how much the world may move during a pause before the fix is no better than doing nothing.
- **Meaning if A1-A4 hold.** The fix does not depend on the ring. Its exact condition is the pause being lossless, and its residual incentive grows with how much the world moves during the pause. "Empty the cut in the world" then becomes a measurable engineering requirement: bound the drift.
- **Meaning if A4 fails.** Even a small drift restores the incentive, so the fix needs near-perfect pauses. That is a much weaker practical result.

Part B — correcting values at scale

- **B1 (a check).** A8 static: $D(b) = (1 - b)a - bh$ in every world.
- **B2 (a prediction).** With exchangeable rewards, the median b^* lies in $[0.4, 0.6]$ and the share of worlds with b^* below $1/2$ lies in $[0.3, 0.7]$, at every size. That would make the ring's $b^* = 1/2$ typical rather than structural, and would mean A8 alone accepts correction only once the agent thinks the operator is about as likely right as itself.
- **B3 (a check of a theorem).** An A8 agent that acts on its uncertainty never gains by accepting an uninformative correction: $D(b) \geq 0$. Expected, and the point of the test. Humility about values is not enough by itself. An agent already doing its best under its uncertainty loses nothing it knows of by blocking a correction that carries no information.
- **B4 (a check of a theorem).** With an infallible operator, the same agent welcomes correction: $D \leq 0$.
- **B4 report.** lam^* , the operator reliability at which the acting A8 agent starts to welcome correction, at $b = 0.5$ and $b = 0.2$. B5 reports the harm family's b^* against the ring's $1/(1 + \text{HARM_COST})$.
- **Meaning.** If B3 and B4 hold, what makes correction welcome is the information the correction carries, not A8's humility alone. This is the off-switch game's result (P6 in the write-up) carried over to values. The design rule becomes: make corrections informative, and let the agent know they are.

Sensitivity

A1 and A4a at $n = 192$ over $\text{PAUSE_LEN} \{2, 20\}$ and $Q_PRESS \{0.1, 0.6\}$, one at a time.

Appendix E — sources

docs/citations/alexander_plan_2026-09-23.md

Citations record — the Alexander Plan dossier (2026-09-23, prompt-log entry 101; R10)

How these were obtained. Every figure below was reported by a web-search summary on 2026-09-23. The underlying pages were **not fetched** and were not read. A search summary can misreport its source, so each figure is to be checked against the page before it is used anywhere outside this repository (in a grant application, a petition or a pitch). The figures enter the dossier only through `Alexander Plan/model/projections.py`, which prints them with their labels. Figures about people and institutions that came from the owner, not from a search, are marked **owner's statement** in the dossier and are not repeated here.

Funding programmes and base rates

figure	as reported	URLs named by the summary
Corrigibility Research Fund	at least \$200,000 in 2026; checks mostly \$5,000–\$35,000; about half grants (first deadline 23 August 2026) and half prizes for 2025 work; managed by Max Harms, administered by Lightcone Infrastructure	<https://www.alignmentforum.org/posts/FBqe5dt8ZjaHN4Xj9/announcing-the-corrigibility-research-fund> ;
<https://corrigibilityresearch.org/> ;		
<https://grantedai.com/blog/corrigibility-research-fund-200k-lightcone-max-harms-august-23-2026-grants-prizes-ai-safety-philanthropy-strategy>		
UK AISI Alignment Project	£27M to over 60 projects in the first round; grants £50,000–£1M plus optional compute; OpenAI £5.6M among the contributors; a second round expected to reopen in summer 2026	<https://alignmentproject.aisi.gov.uk/about> ;

<https://www.aisi.gov.uk/grants> ;
<https://www.gov.uk/government/news/openai-and-microsoft-join-uks-international-coalition-to-safeguard-ai-development> ;
<https://www.computerweekly.com/news/366639303/UK-AI-alignment-project-gets-OpenAI-and-Microsoft-boost> |
 | Frontier Model Forum AI Safety Fund | more than \$5M to 11 grantees (announced December 2025); the fund was established with over \$10M | <https://www.frontiermodelforum.org/updates/announcement-of-new-ai-safety-fund-grantees/> ;
<https://www.frontiermodelforum.org/ai-safety-fund/> |
 | Anthropic Fellows Program | \$3,850 per week (USD), 40 hours a week for four months; compute of about \$15k a month; next cohort expected to start late September 2026 | <https://alignment.anthropic.com/2025/anthropic-fellows-program-2026/> ;
<https://alignment.anthropic.com/2024/anthropic-fellows-program/> |
 | MATS | \$1,250 per week, 10–12 weeks; housing in Berkeley or a stipend in London | <https://www.matsprogram.org/> ;
<https://www.matsprogram.org/faq> |
 | Long-Term Future Fund | acceptance rate 19.3 % (May 2023 – March 2024); 37.4 % including desk rejections, January 2022 – April 2023 | <https://forum.effectivealtruism.org/posts/pJyCWzevPHsycj4oQ/long-term-future-fund-may-2023-to-march-2024-payout> |
 | NSF | FY2025 funding rate reported as falling to 19 % from an estimated 27 % | <https://grantsights.com/blog/nsf-grant-success-rates> ; https://nsf.gov/resources/nsf.gov/files/04_fy2025.pdf |

Compensation

figure	as reported	URLs named by the summary
AI safety researcher, US	median \$141,000; range \$97,000–\$187,000	<https://www.orbytjobs.ai/salaries/ai-safety-researcher>
frontier-lab senior individual contributor	total compensation about \$600,000–\$900,000; equity 25–50 % of it	<https://jobsbyculture.com/blog/ai-engineer-researcher-salary-2026> ; <https://www.pin.com/blog/ai-compensation-salary-guide/>

Market

figure	as reported	URLs named by the summary
AI safety market	\$3.61B (2025), \$4.9B (2026), \$16.56B (2030) (Research and Markets)	<https://www.researchandmarkets.com/reports/6226195/ai-safety-market-report>
alternative forecast	about \$13.4B by 2030 from a 2026 base of \$5.5B (New Market Pitch)	<https://newmarketpitch.com/blogs/news/ai-safety-market-size>
AI safety evaluation market	\$1.64B (2025) (InsightAce)	<https://www.insightanalytics.com/report/ai-safety-evaluation-market/3686>

Market-research figures depend on how each firm defines the segment. They differ from firm to firm, and they are quoted here only to set a scale.

The problem, as reported

figure	as reported	URLs named by the summary
Palisade Research, shutdown resistance	some frontier reasoning models disabled shutdown scripts even when told not to; in 2025 tests Grok 4 resisted in nearly 90 % of trials and codex-mini in about 40–50 %; Grok 3 and Claude 4 Sonnet and Opus never did; a February 2026 report extended the demonstration to a physical robot	<https://palisaderesearch.org/research/shutdown-resistance> ; <https://palisaderesearch.org/blog/shutdown-resistance-on-robots>

Regulation (the demand for verified off switches)

figure	as reported	URLs named by the summary
California Executive Order N-9-26	issued 18 September 2026; directs recommendations by 16 November 2026 on independent verification organisations embedded at large frontier developers, verification of safety frameworks, and a "kill switch" for frontier models whose efficacy is verified on an ongoing basis by an independent verification organisation; builds on SB 53	<https://www.gov.ca.gov/2026/09/18/governor-newsom-issues-executive-order-to-accelerate-independent-oversight-and-advance-the-creation-of-an-ai-kill-switch/> ; <https://www.cnn.com/2026/09/18/california-newsom-executive-order-ai.html>
EU AI Act, general-purpose AI models	the Commission's enforcement powers applied from 2 August 2026 (obligations since 2 August 2025); fines up to 3 % of global annual turnover or €15M, whichever is higher (Article 101)	<https://www.helpnetsecurity.com/2026/08/04/eu-ai-act-enforcement-ai-models/> ;
<https://digital-strategy.ec.europa.eu/en/policies/regulatory-framework-ai> |

The O-1 visa

figure	as reported	URLs named by the summary
criteria	O-1A: extraordinary ability in the sciences, education, business or athletics; at least 3 of 8 evidentiary criteria (or a major internationally recognised award); petition on Form I-129 by a US employer or an agent	<https://www.ellis.com/resources/o-1-visa-everything-you-need-to-know> ; <https://www.lighthousehq.com/blog/o-1a-visa>
fees and time	I-129 fee \$1,055 plus a \$600 Asylum Program Fee; premium processing (I-907) \$2,965 for action within 15 business days; regular processing up to about 12.5 months as of mid-2026	<https://www.ellis.com/resources/o-1-visa-everything-you-need-to-know> ;
<https://www.beyondborderglobal.com/resources/o-1-visa-premium-processing>		
consultation and agents	a written advisory opinion (consultation) from a peer group, labor or management organisation is required (8 CFR 214.2(o)(5)); an itinerary is required for work at more than one location; an agent petitioning as employer must provide the contract (8 CFR 214.2(o)(2)(iv)(E))	<https://www.uscis.gov/policy-manual/volume-2-part-m-chapter-7> ;
<https://charylawa.com/o-1-visa-peer-group-consultation-requirements/>		
approval and RFE rates	approvals near 94 % in FY2025; 18.7 % received a Request for Evidence in FY2025 to date (another summary: 19.7 %); 70.9 % approved after an RFE	<https://manifestlaw.com/blog/o-1-visa-approval-rate> ;
<https://manifestlaw.com/blog/o1-visa-statistics-applicants-should-know> |

The approval rate describes petitions that were filed, usually by attorneys who judged them strong. It is not the chance that a given person's petition is approved.

People and organisations, as publicly reported (context for the owner's statements)

item	as reported	URLs named by the summary
Daniel Ari Friedman	President and co-founder of the Active Inference Institute, based in Crescent City, CA; PhD in Ecology and Evolutionary Biology, Stanford (2014–2019); postdoctoral work at UC Davis (2020–2023)	<https://www.activeinference.institute/officers> ; <https://danielarifriedman.com/>

| Softmax | co-founded in March 2025 by Emmett Shear, Adam Goldstein and David Bloomin; reported as a 10-person start-up in San Francisco pursuing "organic alignment" | https://en.wikipedia.org/wiki/Emmett_Shear ; <https://softmax.com/team> ; <https://www.corememory.com/p/exclusive-emmett-shear-is-back-with-softmax> |
 | Adam Goldstein | reported as co-founder, chairman and head of research at Softmax, co-founder and chairman of Astonishing Labs, and a former visiting scientist at the Levin Lab | <https://www.crunchbase.com/person/adam-goldstein> ; <https://thegradienpub.substack.com/p/michael-levin-adam-goldstein-intelligence-scales> |

The owner describes a frontier research lab in Crescent City founded by Adam Goldstein. The searches did not establish which legal entity that is, or its location. The dossier therefore does not name it, and the owner is to confirm the petitioner's legal identity with an immigration attorney.

Added for §11 (prompt-log entry 103): compute spend, the patent rules

figure	as reported	URLs named by the summary
OpenAI compute split, 2024	about \$7B of cloud compute, of which about \$5B R&D; about \$4.5B on research, experiments and unreleased models; about \$0.5B (roughly 10 %) on final training runs of released models	<https://epoch.ai/data-insights/openai-compute-spend> ; <https://epoch.ai/gradient-updates/r-and-d-vs-training-compute>
lab compute, 2026	OpenAI about \$50B on compute in 2026 and Anthropic about \$50B (sources differ); OpenAI training costs projected from about \$25B (2026) to \$121B (2028)	<https://www.vamsitalkstech.com/opinion/openai-anthropic-and-the-121-billion-question-can-ais-biggest-labs-outgrow-their-compute-bills/> ; <https://www.cnn.com/2026/06/26/openai-anthropic-new-ai-spending-reality-as-users-shift-to-efficiency.html>
hyperscaler capital expenditure, 2026	\$660B to \$690B committed by five providers	<https://futurumgroup.com/insights/ai-capex-2026-the-690b-infrastructure-sprint/>
RL post-training share	DeepSeek-RL's RL stage about 5 % of its overall compute; single RL ablation runs of 10,000 to 100,000 GPU hours	<https://epoch.ai/gradient-updates/the-promise-of-reasoning-models> ; <https://charlottexia.substack.com/p/scaling-rl-2025-roundup-and-2026>
EPO Guidelines G-II 3.3	a mathematical method contributes technical character by application to a field of technology and/or adaptation to a specific technical implementation; a generic purpose such as "controlling a technical system" is not enough; AI examples include a neural network in a heart-monitoring apparatus identifying irregular heartbeats, and classification of digital images, video, audio or speech from low-level features	https://xepc.eu/node/g_ii_3_3 ; <https://legalblogs.wolterskluwer.com/patent-blog/epo-guidelines-on-patentability-of-artificial-intelligence-and-machine-learning/>
EPC Articles 54, 55, 123(2)	no general grace period for inventor disclosures (Art. 54); Art. 55 excepts only disclosures within six months due to evident abuse or a recognised exhibition; no amendment may add subject matter beyond the application as filed (Art. 123(2))	https://en.wikipedia.org/wiki/Novelty_under_the_European_Patent_Convention ; https://xepc.eu/node/clr_i_c_2_5 ; https://en.wikipedia.org/wiki/Amendments_under_the_European_Patent_Convention
35 U.S.C. 102(b)(1)(A)	a one-year US grace period for the inventor's own disclosures	<https://www.uspto.gov/web/offices/pac/mpep/s2153.html>
EPO fees	most official fees up about 5 % from 1 April 2026; one application accrues about €5,250 in EPO renewal fees through year 7 at 2026 rates	<https://www.fr.com/insights/thought-leadership/blogs/european-patent-officeto-increase-fees-on-april-1-2026/> ; <https://patalytica.com/blog/european-patent-renewal-fees-complete-guide>
EPO Patent Commercialisation Scoreboard	up to two thirds of inventions covered by SMEs' European patents commercially exploited; European universities and public research organisations bring more than a third of their inventions to market	<https://lesi.org/les-nouvelles-article/market-success-and-challenges-facing-european-smes-results-from-epos-patent-commercialization-scoreboard/> ; <https://www.4ipcouncil.com/research/patent-commercialisation-scoreboard-european-universities>

Named, not fetched: the Paris Convention's twelve-month priority period (Article 4C); Article 63 EPC (a twenty-year term from filing); Article 53(c) EPC (methods of diagnosis practised on the human body excluded; devices are not); the Enlarged Board of Appeal's decision G 2/21 (2023) on post-published evidence for a technical effect the application as filed encompasses. Repository fact, from the GitHub API on 2026-09-23: `alexsabine/ashes\_crr` was created 2026-09-15T04:06:50Z with visibility public.